

**Advanced Parallel Computing**  
**Summer term 2019**

## Exercise 3

- Return electronically until Tuesday, May 14, 9:00 am
- Include name on the top sheet.
- A maximum of two students is allowed to work jointly on the exercises.

### 3.1 Reading

Read the following two papers and provide reviews as explained in the first lecture (see slides):

- David A. Wood and Mark D. Hill. 1995. Cost-Effective Parallel Computing. *Computer* 28, 2 (February 1995), 69-72. DOI=<http://dx.doi.org/10.1109/2.348002>
- P. Stenstrom. A survey of cache coherence schemes for multiprocessors. *Computer*, vol. 23, no. 6, pp. 12-24, June 1990. doi: 10.1109/2.55497

(25 points)

### 3.2 Shared counter

Implement a Pthread program which performs the following calculation: a counter (initialized to zero) has to be incremented **C** times, so that at the end of the program it has the value **C**. A configurable number of threads (**N** threads in total) shall cooperatively increment the counter. Use a mutex to ensure mutual exclusion when incrementing the counter. Before the threads start incrementing the counter, all threads have to synchronize using a barrier. This avoids the visibility of start-up effects and maximizes contention. **N** and **C** should be command line parameters of your program.

In pseudo-code, each thread should do the following:

```
Barrier();  
for (i = 0 to (C/N)-1) {  
    Mutex_Lock();  
    inc(counter);  
    Mutex_Unlock();  
}
```

Develop a suitable program, first without the mutex. Depending on a sufficient large **C**, data races should result in a mismatch between counter and **C** at the end of the program. Reproduce this behavior. Provide an interpretation why this is dependent on the value of **C**.

Now, implement the program using a mutex to protect the critical section. For a varying **C**, ensure that after execution the counter matches **C**, i.e. there are no race conditions anymore.

Conduct your experiments on **moore**.

(10 points)

### 3.3 Shared counter revisited

Start with the program from the previous exercise. Implement two alternative counter update methods and answer the questions:

1. Use an atomic operation to increment the counter. Obviously, now no locking is required anymore (why?). For gcc, the following intrinsic can be used:

```
uint32_t var = 0;
__sync_add_and_fetch(&var, 1);
```

In this example, variable `var` is atomically incremented. Feel free to use any other atomic intrinsic.

2. Now, use the atomic operation (or another suitable one) to implement your own locking mechanism. I.e., implement a `lock_rmw(*lock)` and an `unlock_rmw(*lock)` function, by relying on atomic operations. Why are atomic operations required in this case?

Develop your programs and perform initial testing on `moore`. Validate the correctness of these two new programs with the same methodology as in exercise 3.2 (for a varying **C** and **N**, ensure that after execution the counter matches **C**, i.e. there are no race conditions anymore).

(25 points)

### 3.4 Shared counter performance analysis

Now, measure the overall execution time using suitable functions (e.g. `clock_gettime` or `gettimeofday`). Report the overall execution time and the derived number of updates per second for a varying number of threads (1-48) and sufficiently large number of updates (providing stable results). For this experiment, use the computer `moore` (48 cores in total). As `moore` is often heavily used, please ensure that you only use it for performance experiments.

	PTHREAD MUTEX		ATOMIC INCREMENT		LOCK_RMW	
Thread count	Execution time	Updates per second	Execution time	Updates per second	Execution time	Updates per second
1						
2						
4						
8						
12						
16						
24						
32						
40						
48						

Run this experiment with all four implementations. Report results in the table above, include appropriate units (ns, us, ms, s). Include a graphical representation here (varying number of threads on x-axis with updates per second on y-axis). Interpret your results.

(15 points)

**Total: 75 points**