

Advanced Parallel Computing

11.06.2019

Students:

Jona Neef

Nikolas Krätzschar

Philipp Walz

Exercise 6

6.1 Reading

Calin Cascaval 2008. Software Transactional Memory: Why Is It Only a Research Toy?

In their paper Călin Cașcaval and his colleagues present a new software transactional memory framework from IBM and compare it with the Intel STM and Sun TL2 STM. They also discuss the various problems involved in programming with TM, including both HTM and STM.

The main finding was that performance was very poor in all STM implementations. While with the k-means algorithm the IBM STM implementation achieved a single threaded performance with at least 4 cores, this was not possible with vacations for any of the implementations, not even with 8 cores. The developers explain this by the enormous overhead of the TM and analyze it with regard to the different STM operations. Especially read and write barriers do not perform well, which is the reason why further research should try to minimize them.

I accept the opinion of the authors that the enormous complexity generated by TM programming (whether in hardware or software) ultimately reduces productivity and makes it difficult to use it outside of research work.

Aleksandar Dragojevic 2011. Why STM can be more than a research toy.

In their paper "Why STM can be more than a research toy" Dragojevic et. al. show that in contrast to the findings of Cascaval et al. transactional memory (TM) does have its reason for existence, if used in a hybrid software-hardware implementation, and can often outperform sequential, nontransactional code.

To evaluate TM's viability based on their level of overhead and level of speedup. Their key insight is that parallel applications exhibiting high contention are not the primary target for STM therefore they shouldn't be used as a benchmark.


Secondly the authors show that features like compiler instrumentation and explicit, nontransparent privatization leverage transactional memory to a degree that outperforms sequential code in every but one tested workload.

We strongly agree with the authors work. In our opinion it is always a matter of where to set the benchmark. For sure transactional memory has severe downsides compared with other advanced parallel mechanisms. However for the typical programmer without special knowledge of parallel programming, TM delivers easy-to-use interface that could be build into a wide range of future commodity software.

6.3 Parallel Prefix Sum Analysis

```
1 thread_cnt,default,membind=1,interleaved=4
2 1,2680.48,2634.48,3071.54
3 2,1869.26,1791.61,2054.74
4 4,1534.49,1622.36,1840.58
5 8,1290.23,1306.72,1604.86
6 12,1576.24,1688.25,1712.32
7 16,1222.73,1308.41,1533.75
8 24,1317.87,1436.56,1481.54
9 32,1222.38,1196.25,1398.56
10 40,1267.12,1418.98,1451.64
11 48,1212.6,1275,1366.95
```

Performance per parallel threads for prefix sum calculation



The chart shows the overall prefix sum calculation performance per different thread counts (t).

We see, that using two threads instead of sequential execution yields the highest relative speedup.

Furthermore, at the points $t = [12, 24, 40]$ we can see that the performance drops. This is due to the fact that the thread count is not a power of two as with the other values.

Surprisingly, the interleaved numa configuration does not outperform the default or membind configuration. One explanation for this is, that although our parallel program accesses the array elements in an interleaved fashion (i.e. in the first run thread 0 accesses 0 and 1, thread 1 accesses 2 and 3, and so on), the `interleaved=4` configuration distributes the array in 4 whole slices to the different nodes. Clearly this subdivision does not bring any advantage to our algorithm. To further optimize performance to our algorithm we must distribute the array in a more granular way.