

# ADVANCED PARALLEL COMPUTING

## LECTURE 01 - INTRODUCTION

Holger Fröning  
[holger.froening@ziti.uni-heidelberg.de](mailto:holger.froening@ziti.uni-heidelberg.de)  
Institute of Computer Engineering  
Ruprecht-Karls University of Heidelberg

# ABOUT ME & THIS COURSE



**NVIDIA®**



Professor at Heidelberg University, Germany

Currently advising 9 PhDs, 4 being external (3 on ML)

PhD from Mannheim University

Post-Doc, TU Valencia & Heidelberg University (FPGAs, ASICs)

Visiting professor, TU Graz (2015 - first contact with ML)

Visiting scientist, NVIDIA Research (Santa Clara, CA) (2016)

Performance, energy efficiency & programmability for future & emerging technologies

Computer engineers/architects with a focus on low-level software layers

High-performance computing, machine learning & data analytics

Bipolar computing landscape: super small & super big

This course

Basics and advanced methods for scalable multi-core computing systems

Research-driven & forward-looking

# OBJECTIVES & PREREQUISITES

## Objectives: The students ...

- understand advanced concepts and principles of parallel architectures
- will be able to develop optimized programs for parallel architectures
- can use the learned structures to develop new architectures

## Methodology

Lectures focus on current architectures, trends, technology constraints

Exercises: Practical hands-on, reading

Excessive programming & paper reviewing

## Prerequisites

Required: Solid knowledge about computer architecture and C language

Recommended: Course “Introduction to HPC” & “GPU Computing”

# ORGANIZATION

Lectures - 2 hours/week - [holger.froening@ziti.uni-heidelberg.de](mailto:holger.froening@ziti.uni-heidelberg.de)

Tuesday, 14:00 ct

Exercises - 2 hours/week - [bernhard.klein@ziti.uni-heidelberg.de](mailto:bernhard.klein@ziti.uni-heidelberg.de)

Tuesday, 16:00 ct

Groups of up to two students allowed - individual work must be visible

Mixture of reading/exercises/programming/experiments

Second half of term will be project work

One final oral or written exam

Prerequisite: see exercise announcement

6 CPs

# ASSIGNMENTS

Practical exercises: usually coding & experiments

Reading & feedback based on paper review

Ideal review here is 3 sentences for each of the following:

1. Primary contribution
2. Key insight of the contribution
3. Your opinion/reaction to the content

Review: rating relative to all other papers (of this venue)

Strong reject, weak reject, weak accept, accept

Old papers: optionally include some comments on how right this paper was

Project assignment for second half of the term (tba)



# ADDITIONAL READING

## Books

Culler et al., Parallel Computer Architecture: A Hardware/Software Approach, Morgan Kaufmann

Hennessy/Patterson, Computer Architecture: A Quantitative Approach, Morgan Kaufmann


(Keckler et al., Multicore Processors and Systems, Springer)

## Publications/Conferences

ISCA, HDCA, ASPLOS, PACT, IPDPS, ICPP, ISPASS, HPDC, ...

See ACM/IEEE websites (or author's web site for limited copies)

Google Scholar and similar



# QUESTIONNAIRE

Pthread & OpenMP

Amdahl's law

Difference between coherence and consistency

Intel MIC

Exponential back-off

Dennard scaling

Victim cache

Data-race free

Cache conflict

# CURRENT COMPUTING TRENDS



# PARALLEL COMPUTING - DEFINITION

A collection of processing elements that communicate and cooperate to solve large problems fast.” - Almasi & Gottlieb, 1989

## Issues

- How large/powerful/scalable?

- Methods for communication/synchronization?

- Interconnection network and operations?

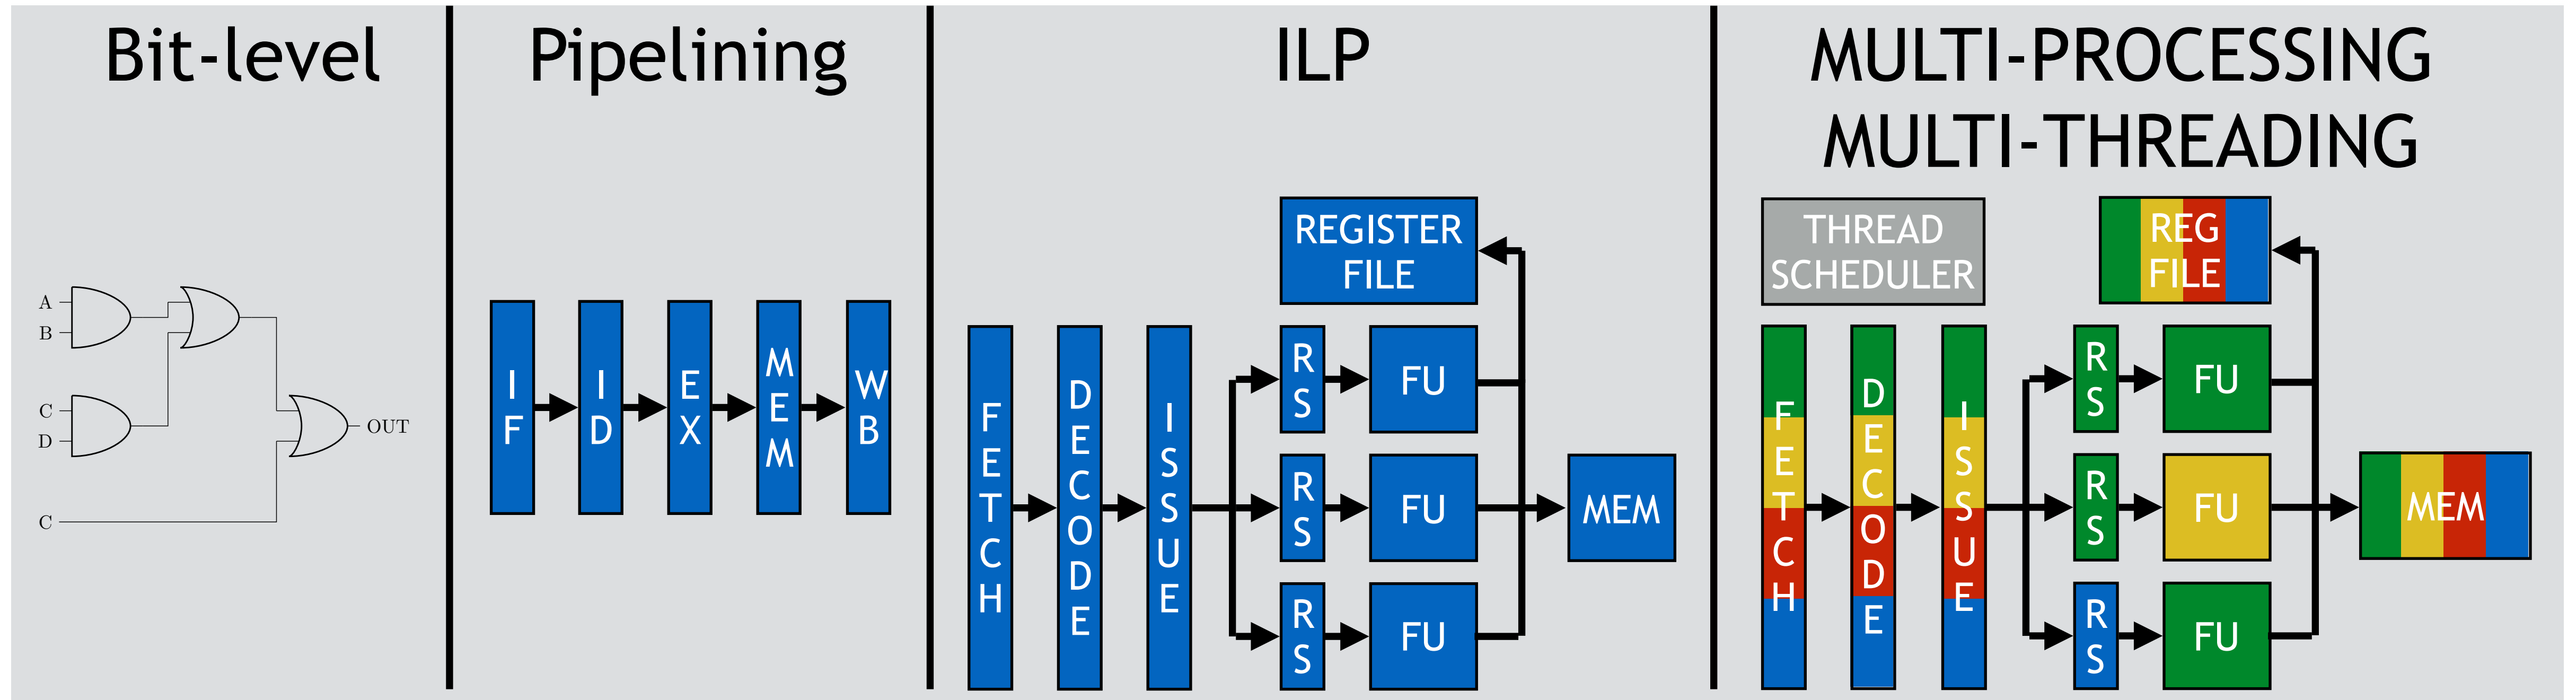
- Primitive abstractions for the programmer?

Note that this definition gears to increase performance

- Other reasons: availability (not covered in this course)

Keys: concurrency and parallelism

# DIFFERENT LEVELS OF PARALLELISM



Main motivation: performance!

Limited amounts of parallelism

Coarse grain parallelism scales better

# DIFFERENT LEVELS OF PARALLELISM

Instruction-Level Parallelism (ILP) ==> Pipelining, Superscalar

Parallelism within one instruction stream

Limited parallelism

Huge amount of dependencies and branches

Thread-Level Parallelism (TLP) ==> SMT, Multi-Core, Cluster

Parallelism in multiple independent instruction streams

Less amount of dependencies, no limitations due to branches

Limited by maximal concurrently executable I-streams

Data-Level Parallelism (DLP) ==> GPUs

Vectorization techniques

Applying one operation on multiple elements of a data structure

Parallelism dependent on data structure

Request-Level Parallelism (RLP) ==> Datacenters & WSC

Huge amount of concurrent requests

Unstructured, randomized patterns

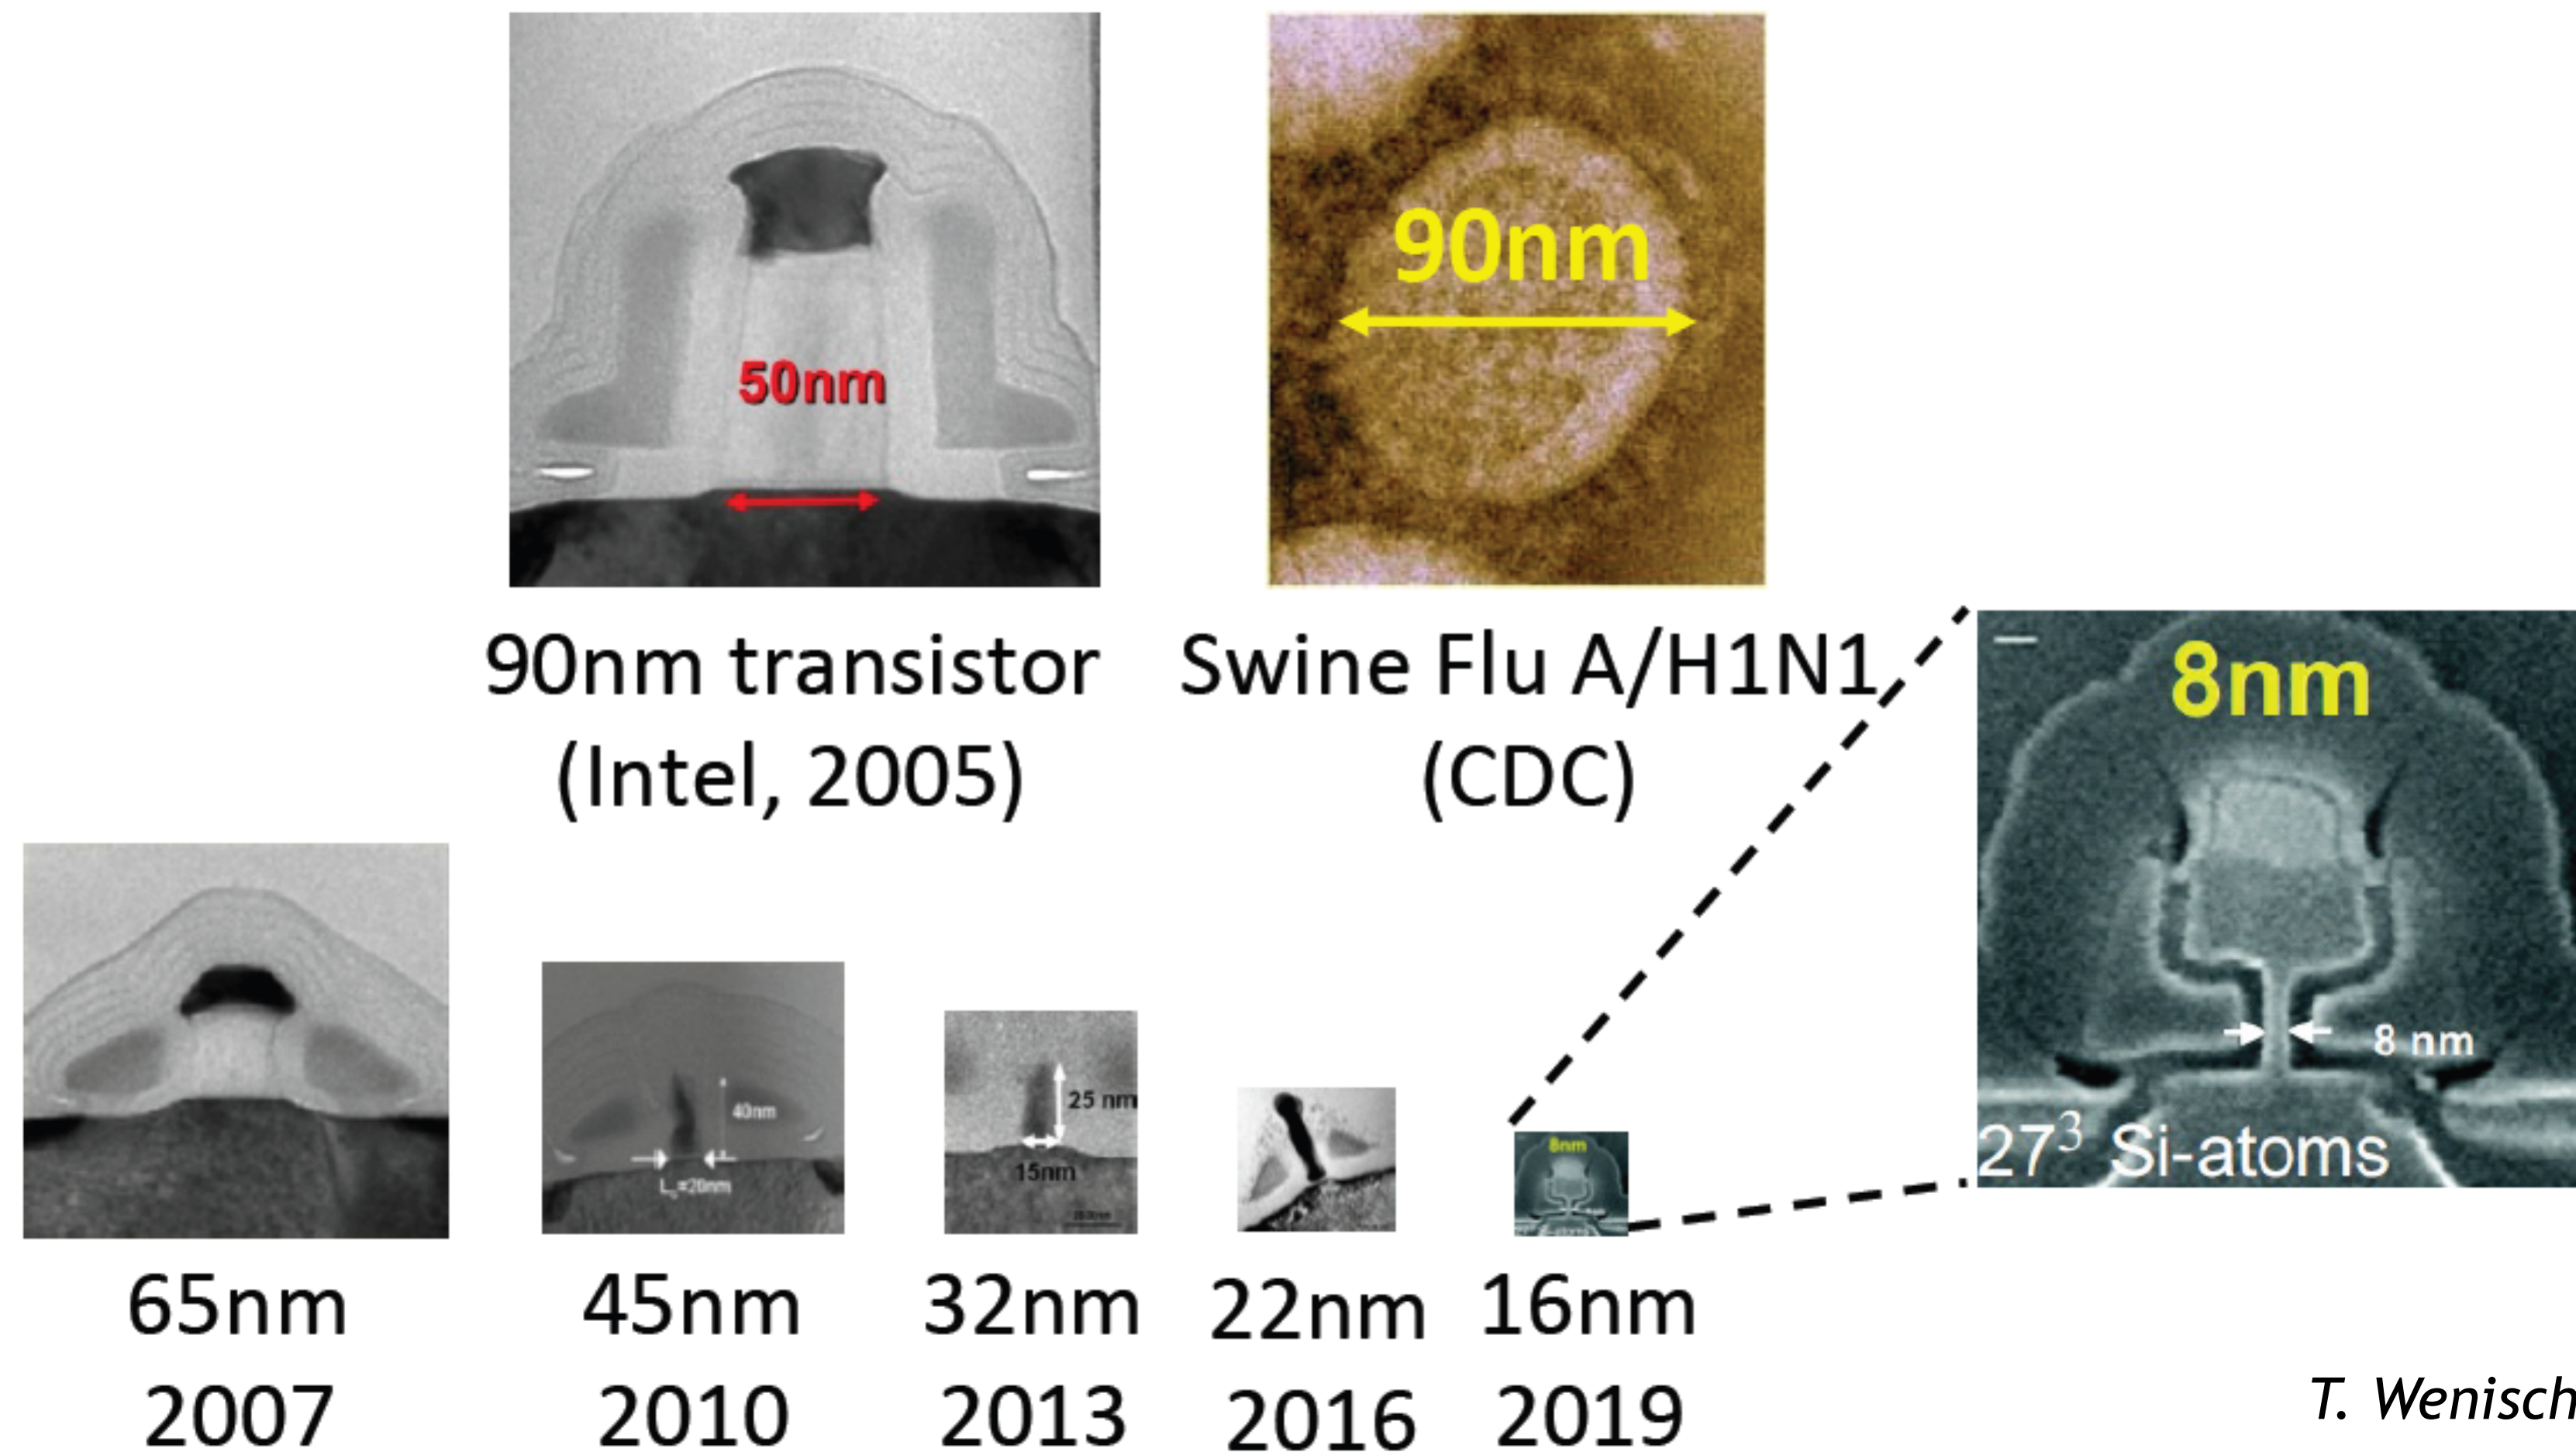
Basics of computer  
architecture

Course “Parallel Computer  
Architecture”, this course

Course “GPU Computing”



# OBSERVATION: TRANSISTOR COUNT



Moore's law will probably continue for another decade

# OBSERVATION: ILP WALL

Applications don't contain enough ILP

IPC for 6-issue is higher, but not 3x in comparison to 2-issue

Main limitations: Stalls due to cache misses and dependencies




Figure 4. IPC Breakdown for a single 2-issue processor.





Figure 5. IPC Breakdown for the 6-issue processor.




# OBSERVATION: LIMITED SIGNAL REACH

Propagation delay increases,  
frequency increases/increased  
=> Signal speed decreases  
=> Limited signal reach





# OBSERVATION: POWER WALL



*K. Yelick*

Inflection point in 2005 - end of Dennard scaling

# THE FUNDAMENTAL TRANSITION TO MULTI-CORE

$$P = afCV^2 + VI_{leakage}$$

a constant,  $f \sim V$

Transition to multi-/many-core

Given a single-core design

Reduce clock to 80%

Power  $\sim f^3 \Rightarrow$  now at 51.2%


Add second core

Assume that single-core performance  $\sim f$

Same power budget (envelope), but 1.6x performance

Transition to asymmetry/heterogeneity

Specialization by massive parallelization



# OVERVIEW OF PARALLEL ARCHITECTURES

# PARALLEL PROGRAMMING MODELS

**Programming model:** language & libraries that define an abstract view of a machine

**Control - expose parallelism**

- How parallelism is created

- How are dependencies solved?

**Data - communication, placement and partitioning**

- Shared/Private

- Accessibility

**Concurrency control - synchronization**

- How to coordinate/orchestrate?

- Atomic operations to ensure mutual exclusion

# ARCHITECTURE FOUNDATIONS

## Naming

How is data communicated?

How are control flows addressed?

## Operations

What operations are allowed on named data?

## Ordering

How can control flows coordinate their activities?


## Performance

Latency and Bandwidth

## Data races / Race condition

A race condition is a flaw in a computing system as the result of an unexpected and critically dependent sequence of events.

# ARCHITECTURE OVERVIEW



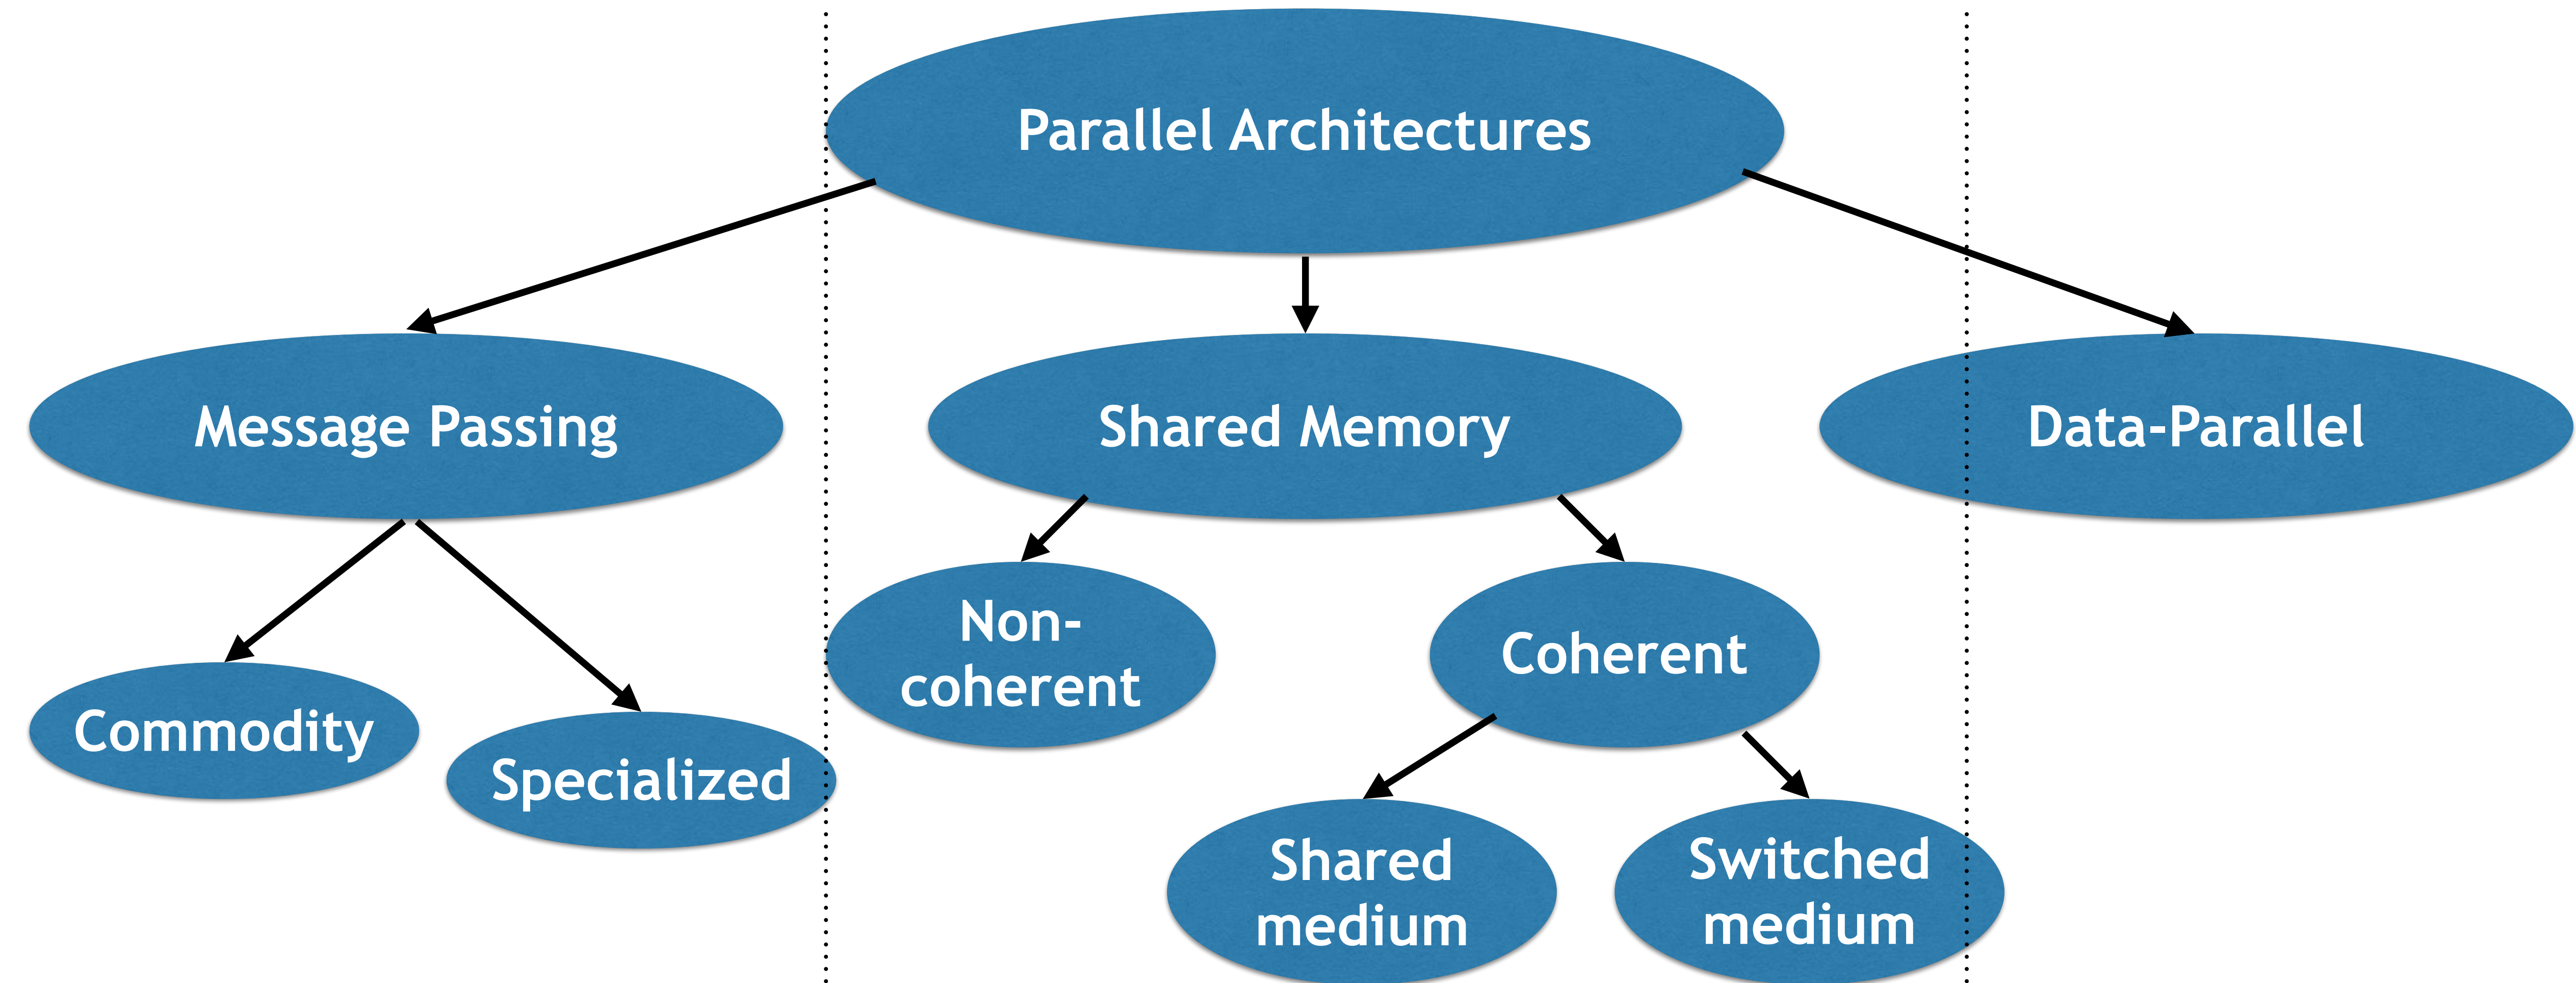
IO: Message passing

Memory: Shared Memory

Processor: SIMD, VLIW, CUDA, dataflow architectures



# ARCHITECTURE TAXONOMY



Course IntroHPC

This course

Course GPUComp<sub>21</sub>

# MULTI-/MANY-CORE TREND

Moore's law gives us plenty of transistors, question is how to use them

- Memory ruled out as DRAM technology is different

- ILP is too limited, no big uniprocessor

- Big caches suffer from increased hit times, increasing average memory access time

- Thus: use them for caches or for processing cores?

Chip Multi-Processors (CMP, aka multi-core)

- Foundation of today's computing arena, making parallel programming pervasive

Data-Parallel Architectures (vector processors, GPUs, MICs)

- Specialized styles of computing for improved performance & energy efficiency

# SCHEDULE

Date	Lecture	Exercise
23.04.19	Intro	Reading
30.04.19	Shared Memory Architectures	Pointer Chasing
07.05.19	Snooping coherence	Shared counter (manual lock), atomics
14.05.19	Synchronization I: locks	MCS lock
21.05.19	Synchronization II: barriers	Barrier
28.05.19	Transactional memory	Parallel Prefix Scan
04.06.19	Scalable coherence	Project work
11.06.19	Advanced coherence techniques	-
18.06.19	Memory consistency	-
25.06.19	Applied consistency	-
02.07.19	ML introduction & optimizations	-
09.07.19	Trends & Future	-
16.07.19	<place holder>	Final presentation
23.07.19	Exam	-

# SUMMARY

This course

## Post-Dennard I: concurrency & specialization in compute (and memory)

Strong NUMA effects, effective tolerance techniques

Implications have yet only marginally been addressed

## Post-Dennard II: communication-centric systems

Communication more expensive than compute

Energy is fundamental, no hiding possible

Strong and even dynamic NUMA effects

## Post-CMOS: specialization & complexity

