

ADVANCED PARALLEL COMPUTING

LECTURE 07 - SCALABLE COHERENCE

Holger Fröning
holger.froening@ziti.uni-heidelberg.de
Institute of Computer Engineering
Ruprecht-Karls University of Heidelberg

Some material by Falsafi, Hardavellas, Nowatzky of EPFL, Northwestern, CMU

ANNOUNCEMENTS

Projects due 09.07.

No lecture 16.07.

SCALABLE CACHE COHERENCE

Snooping: totally ordered interconnection network

Requests: Get-S, Get-M, Put-M

Part 1: bus bandwidth

Replace non-scalable, frequency-limited shared medium with a scalable point-to-point interconnection network

Mesh, tori, BMIN, etc...

Part 2: snooping bandwidth

Observation: most snoops result in no action

I.e., cache line is not shared

Replace non-scalable broadcast protocol with more scalable directory protocol

Broadcasts => multicasts

BASICS OF DIRECTORY COHERENCE PROTOCOLS

DIRECTORY COHERENCE PROTOCOLS

Observation: physical address space is statically partitioned for a multi-processor system

- System with multiple memory controllers

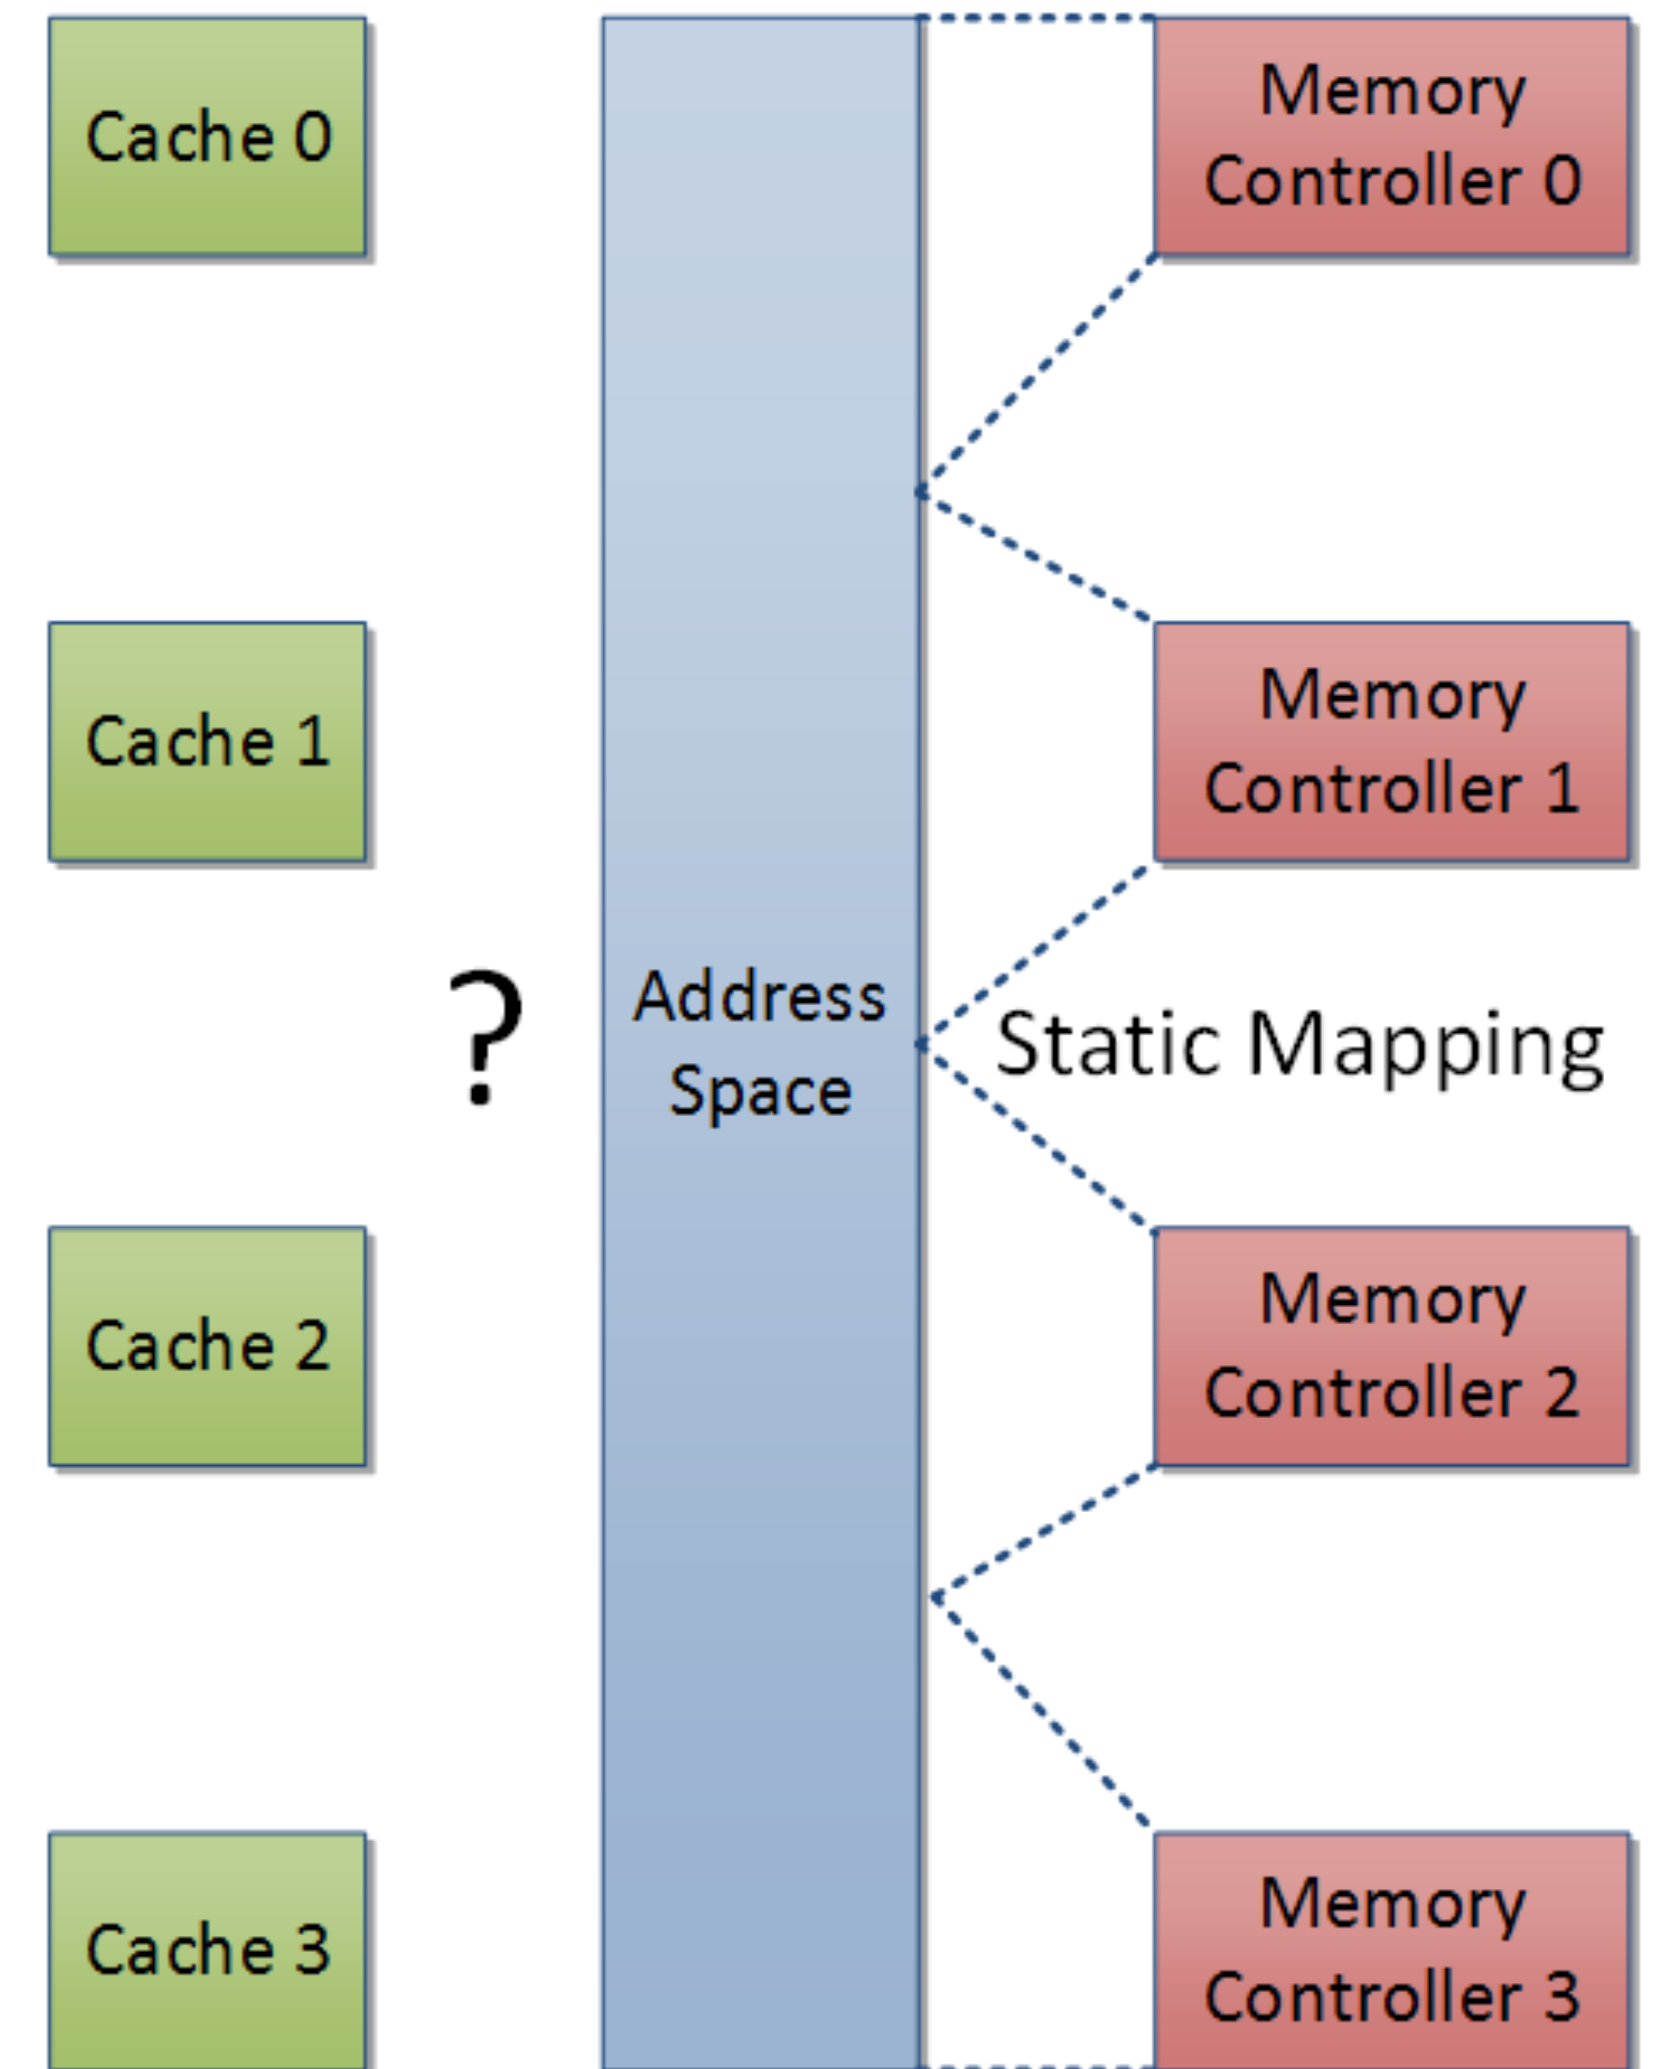
- Easy to determine which memory module “home” holds a given line

Observation: caching is highly dynamic, frequent copying & migration of cache lines

- Difficult to determine which processors have a given line in their caches

- Main reason for bus-based protocols: inherent broadcast probes all caches

- Simple and fast, but non-scalable



DIRECTORY COHERENCE PROTOCOLS

Directories: non-broadcast coherence protocol

Extend memory to track caching information

I.e., references to caches that include copies

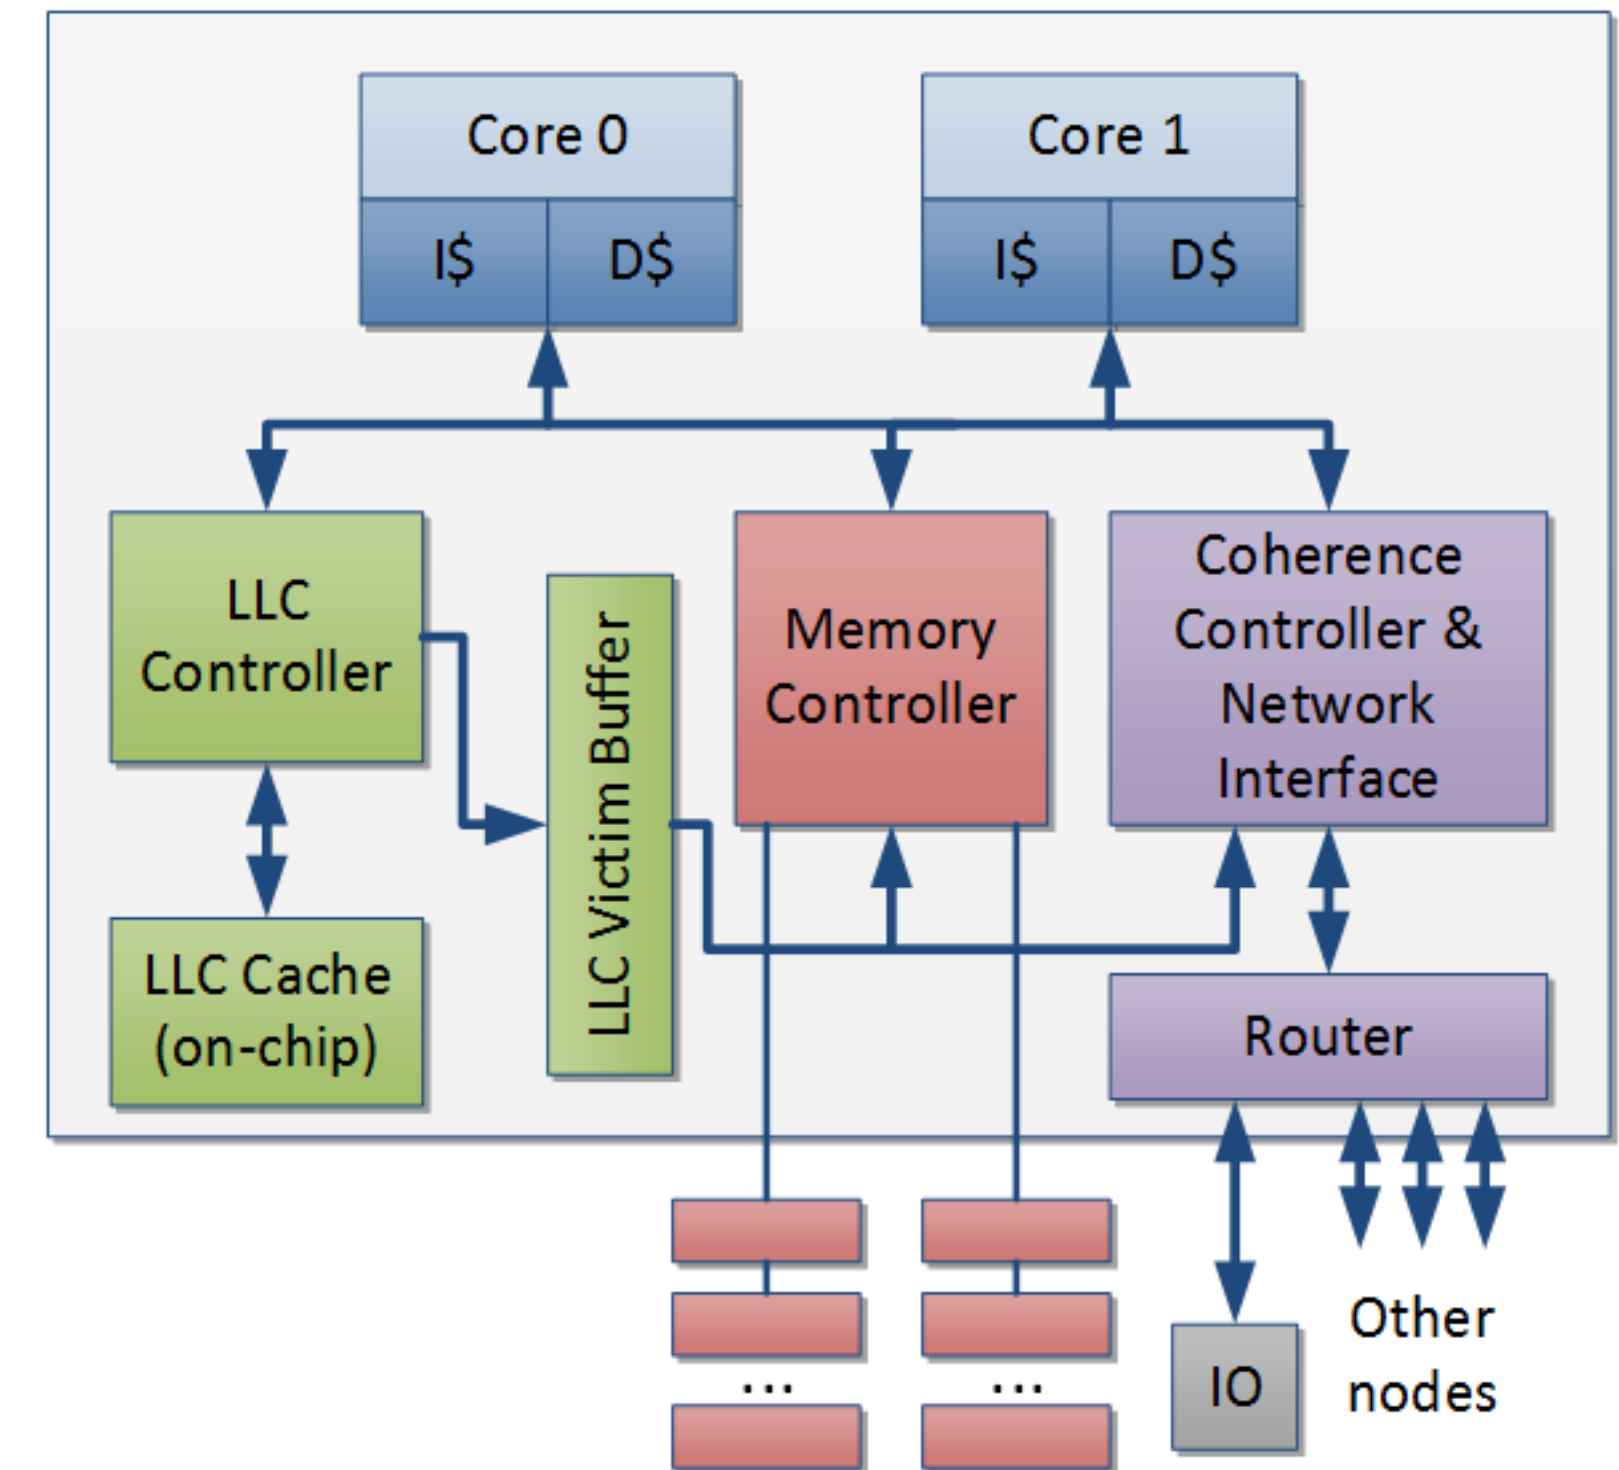
For each physical cache line (that is homed here), track:

Owner: processor that has a dirty copy (modified state)

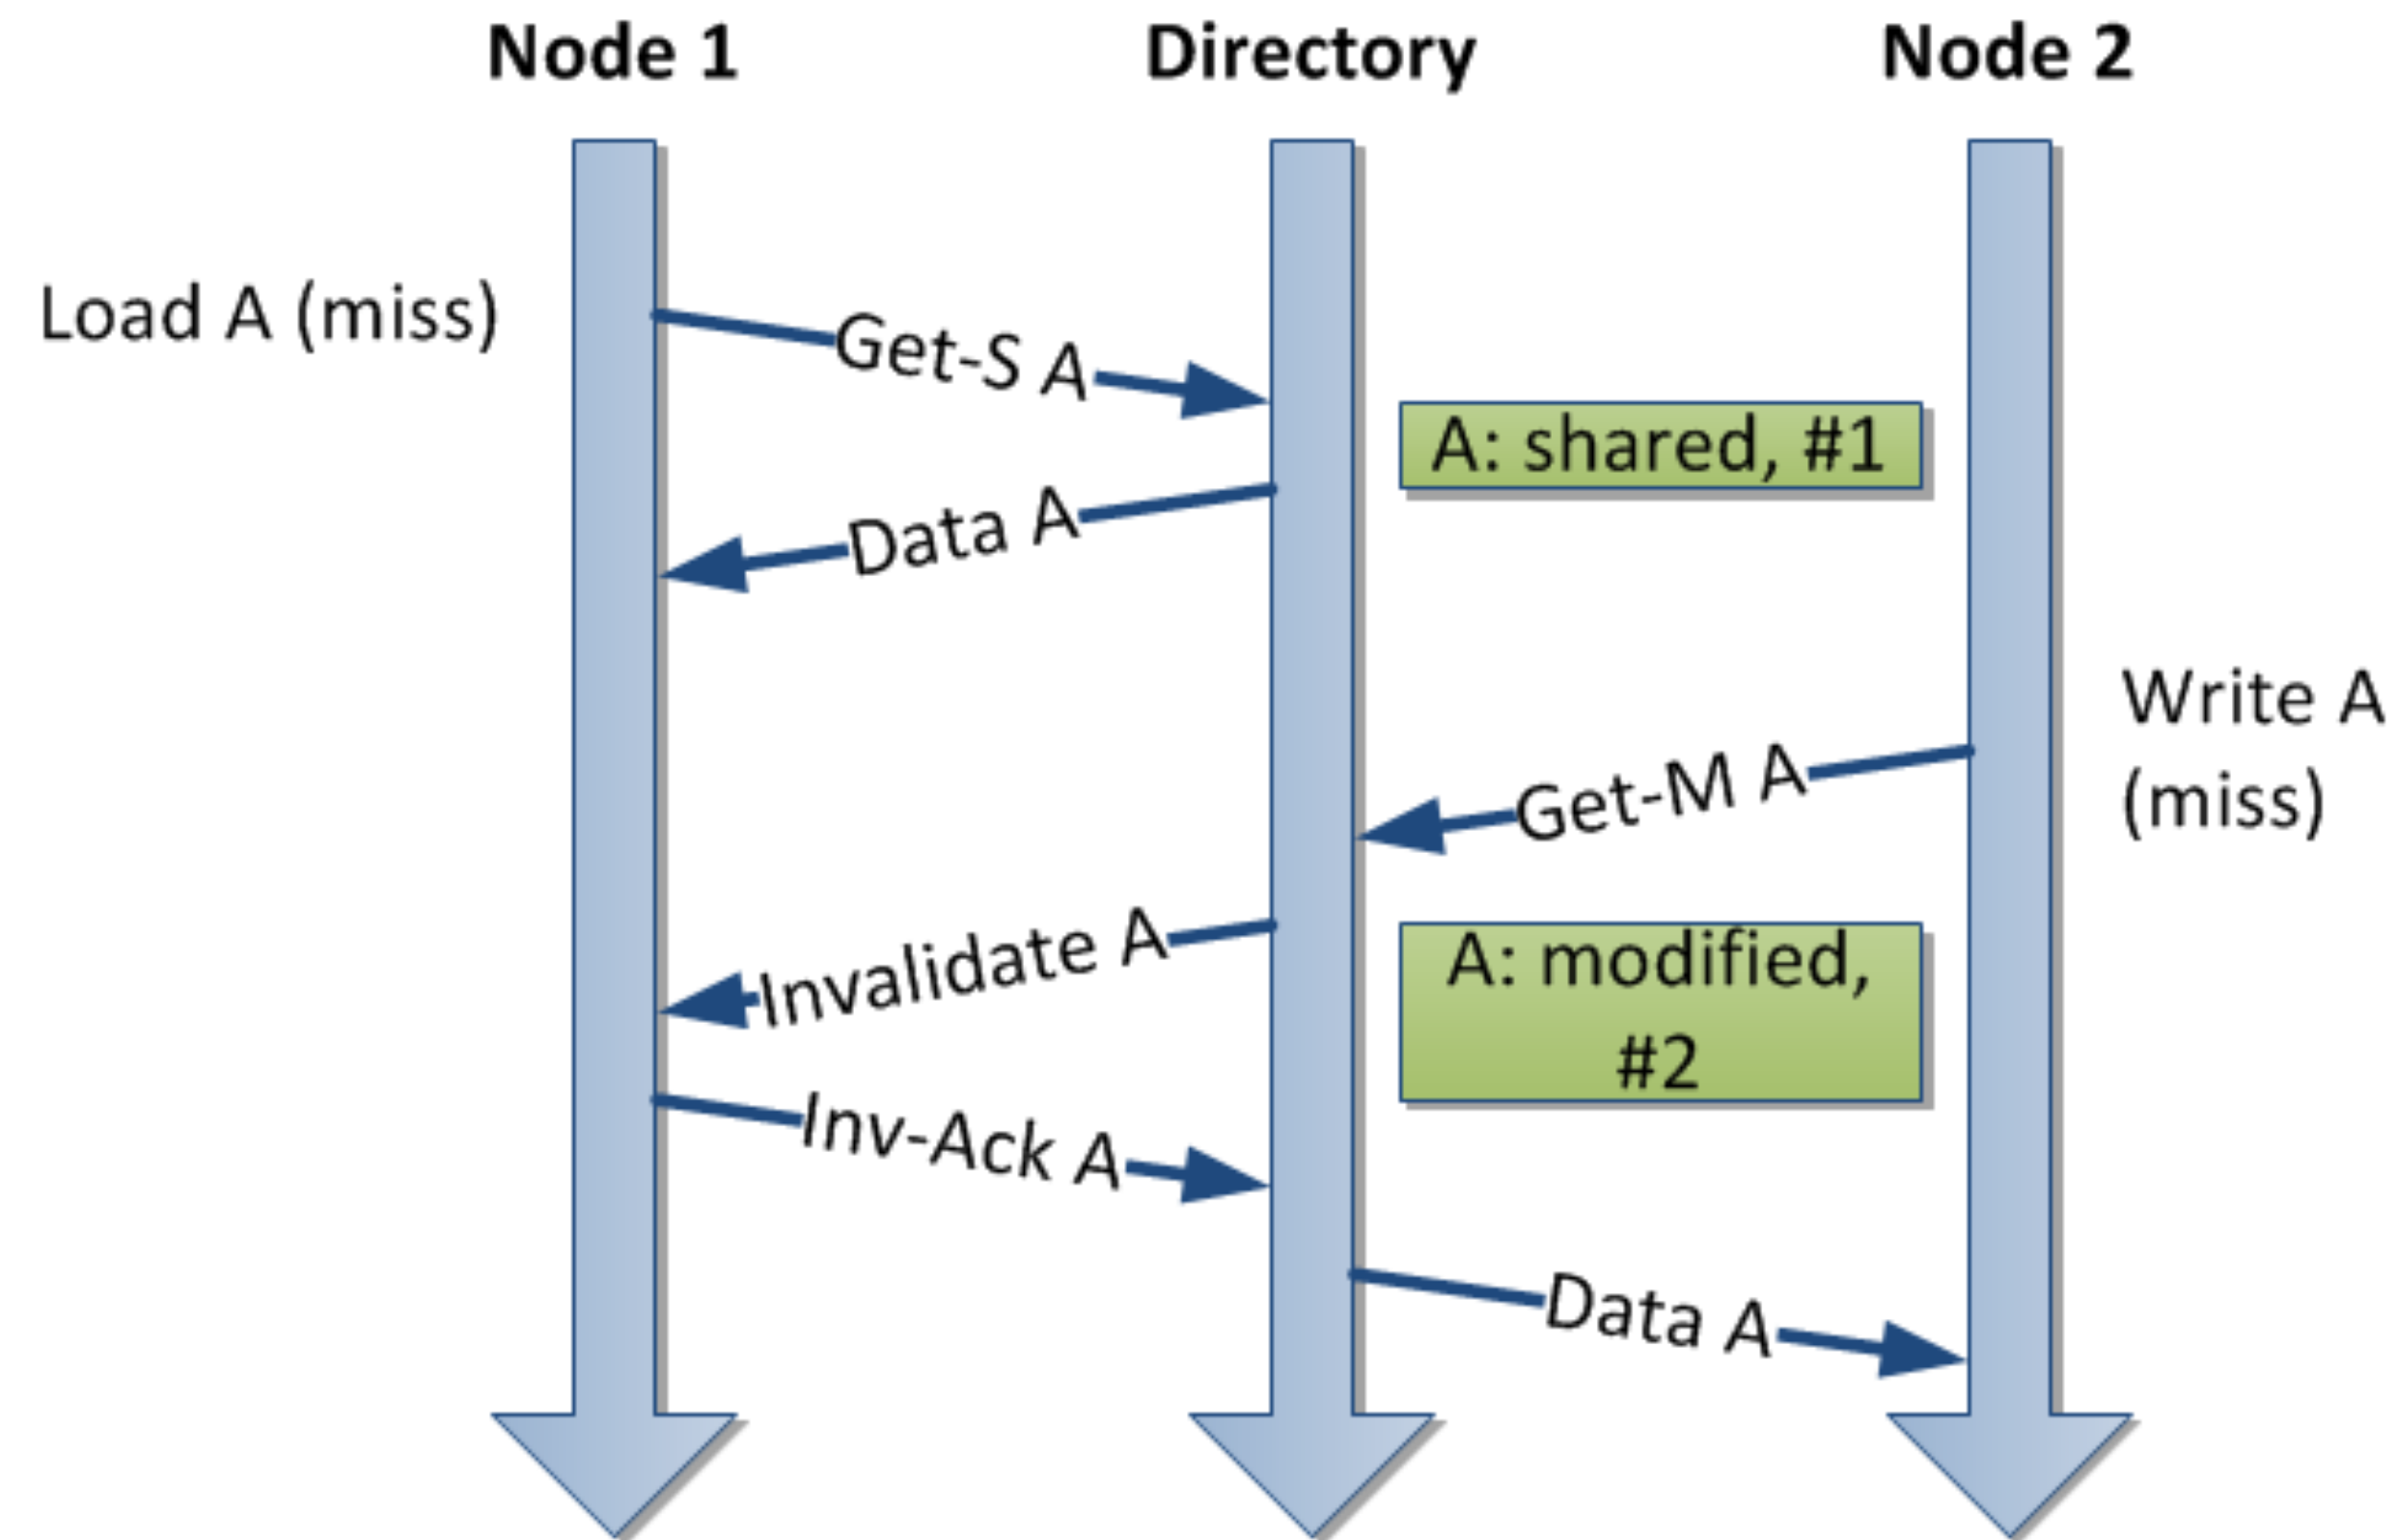
Sharers: processors that have clean copies (shared state)

Each requesting processor sends coherence events to home directory

Home directory forwards events to relevant caches only



BASIC OPERATION: READ/WRITE



Centralized Directory

Distributed Directory

CENTRALIZED DIRECTORY

Single directory contains a copy of the cache tags of all nodes

Advantages

- Send invalidate/update only to the nodes that have copies

- Central serialization point: easy to get strong consistency

- Works like a bus

Disadvantages

- Not scalable, contention point

- Directory size/organization changes with number of nodes

=> Conflictive with scalable point-to-point networks

DISTRIBUTED DIRECTORY

Distribute directory among memory modules

Memory block = coherence block (= cache line)

Home node: node with directory entry

Usually also the main memory storage for cache line


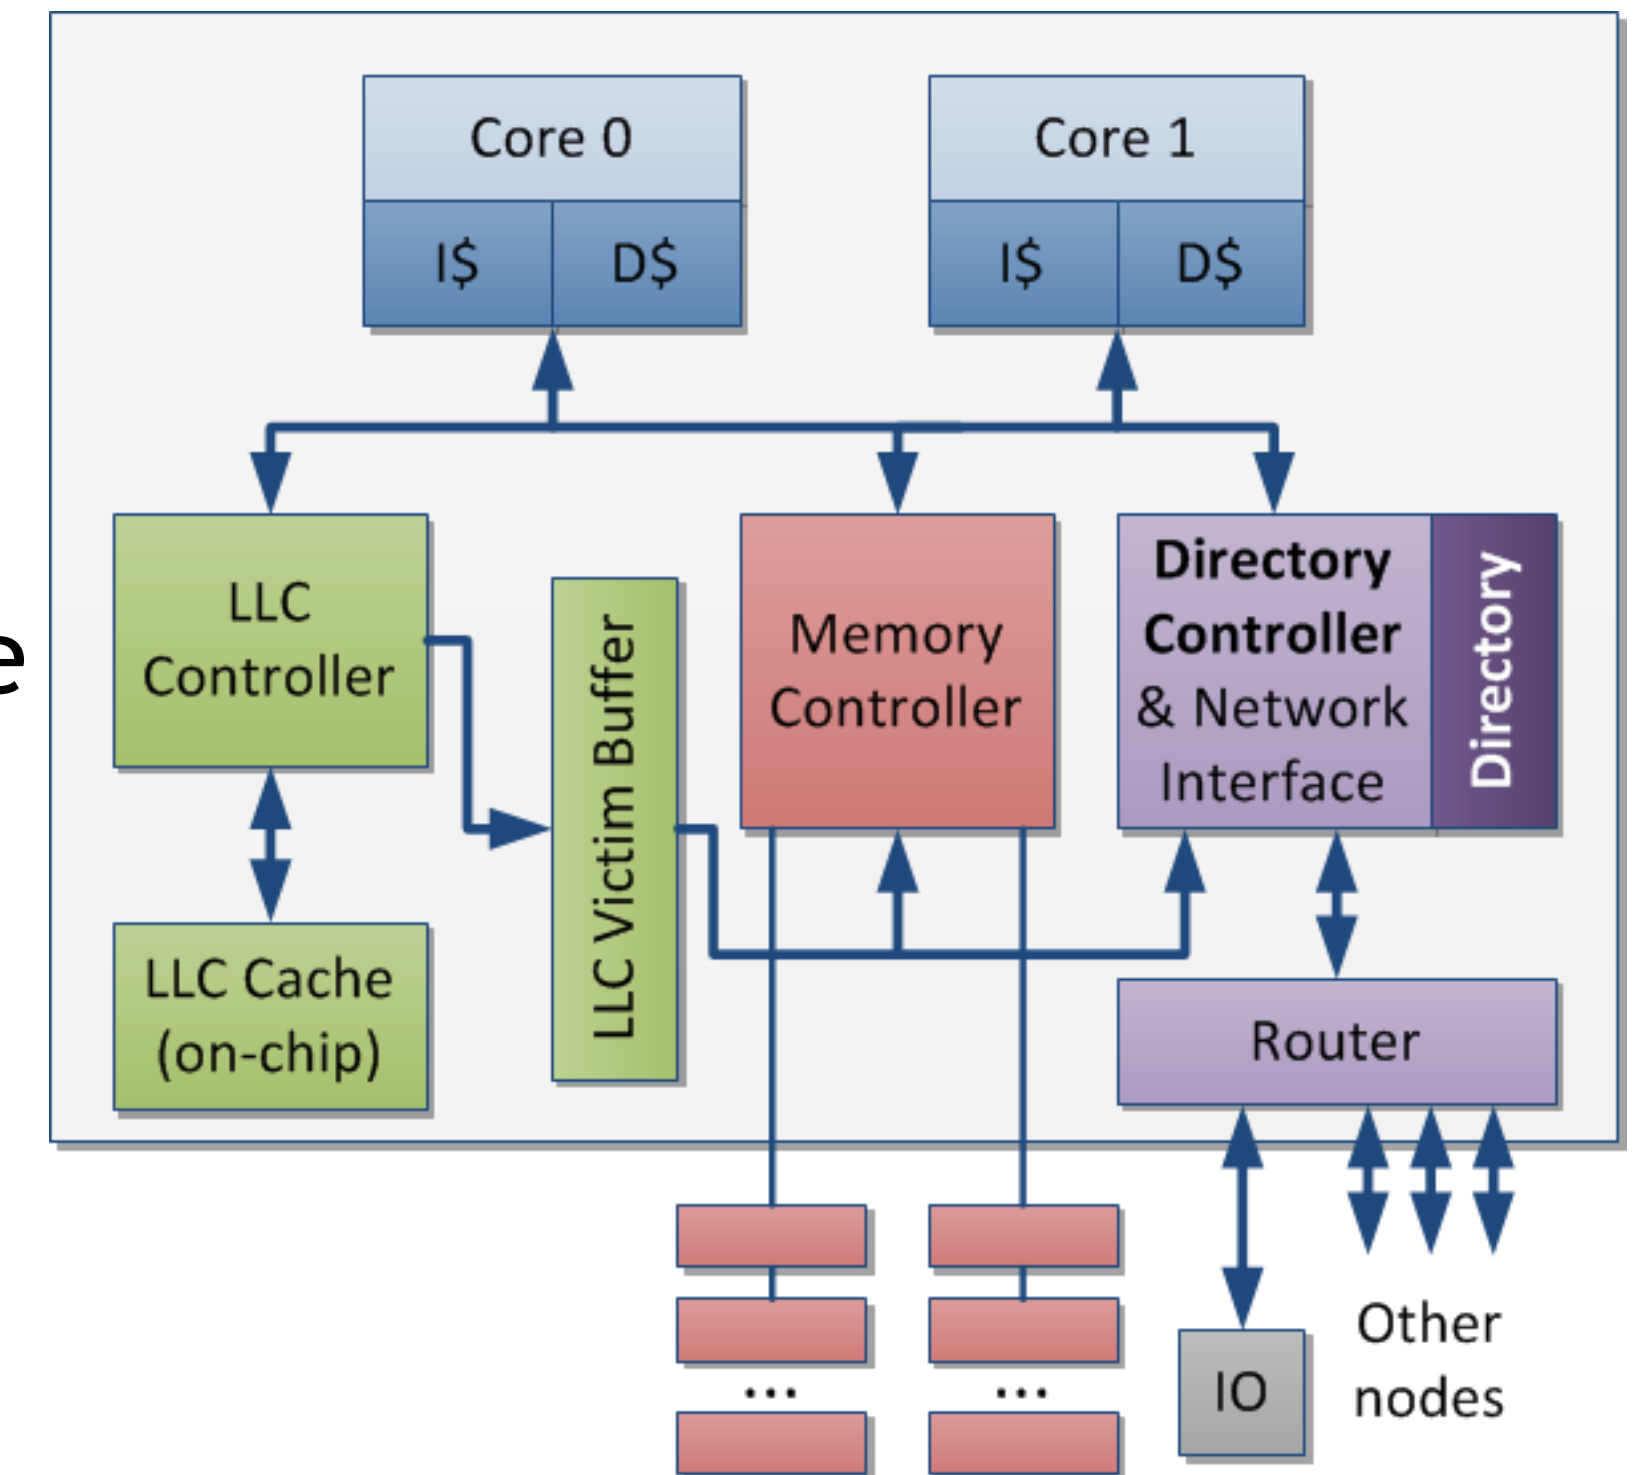
Scalable: directory grows with memory capacity

Common trick: steal bits from ECC for directory state

Problems

Directory can no longer serialize accesses across all addresses

=> Memory consistency becomes responsibility of processor (interface)



CONTENTS OF A DIRECTORY

Directory State

Invalid, Exclusive, Shared, ...

Few stable states

Usually 2-3 bits are enough

Number of outstanding invalidations, ...

Transient states

Pointer to exclusive owner

Sharer list

List of caches that may have a copy

May include local node

Not necessarily precise, but conservative

=> Often 10s of transient states, often with references to node IDs, etc.

Transient state changes frequently, need fast RMW ops

Design options

Keep in directory: scalable (high concurrency), but slow

Keep in separate memory

Keep in directory, use cache to accelerate access

Keep in protocol controller

Transaction State Register File - like MSHRs

One directory entry per memory block (e.g., 64 bytes)

=> Huge directories

=> Entry size matters

CONTENTS OF A DIRECTORY: EXCLUSIVE OWNER

Only one reference => simple node ID (N bits), 2^N nodes

Can share storage with sharers list ==> only 1 bit

Disjoint use

May point to a group of nodes that internally maintain coherence (hierarchy, e.g., via snooping)

Coarse directories (think of multi-core)

May treat local node differently

CONTENTS OF A DIRECTORY: SHARER LIST

Key to scalability - number of sharers can be huge

Must efficiently represent node subsets

Observation: most blocks only cached by 1-2 nodes

However, a few blocks are heavily shared: synchronization variables

Bit vectors: one bit per processor/node/cache

Storage requirements grow with system size -> not scalable

Limited pointer list: fixed number of pointers to node IDs (e.g., 4)

If more than n sharers: (a) recycle one pointer (invalidation), (b) revert to broadcast, (c) handle in software

Linked lists: each node stores pointer to previous and next sharer

Double linked list: Scalable Coherent Interconnect (SCI)

Single linked list: S3.mp

Scalable, but poor performance: long invalidation latency, replacement difficult (in particular for single linked list)

DIRECTORY EXAMPLE

DIRECTORY EXAMPLE

L: Local Node

Node initiating the transaction

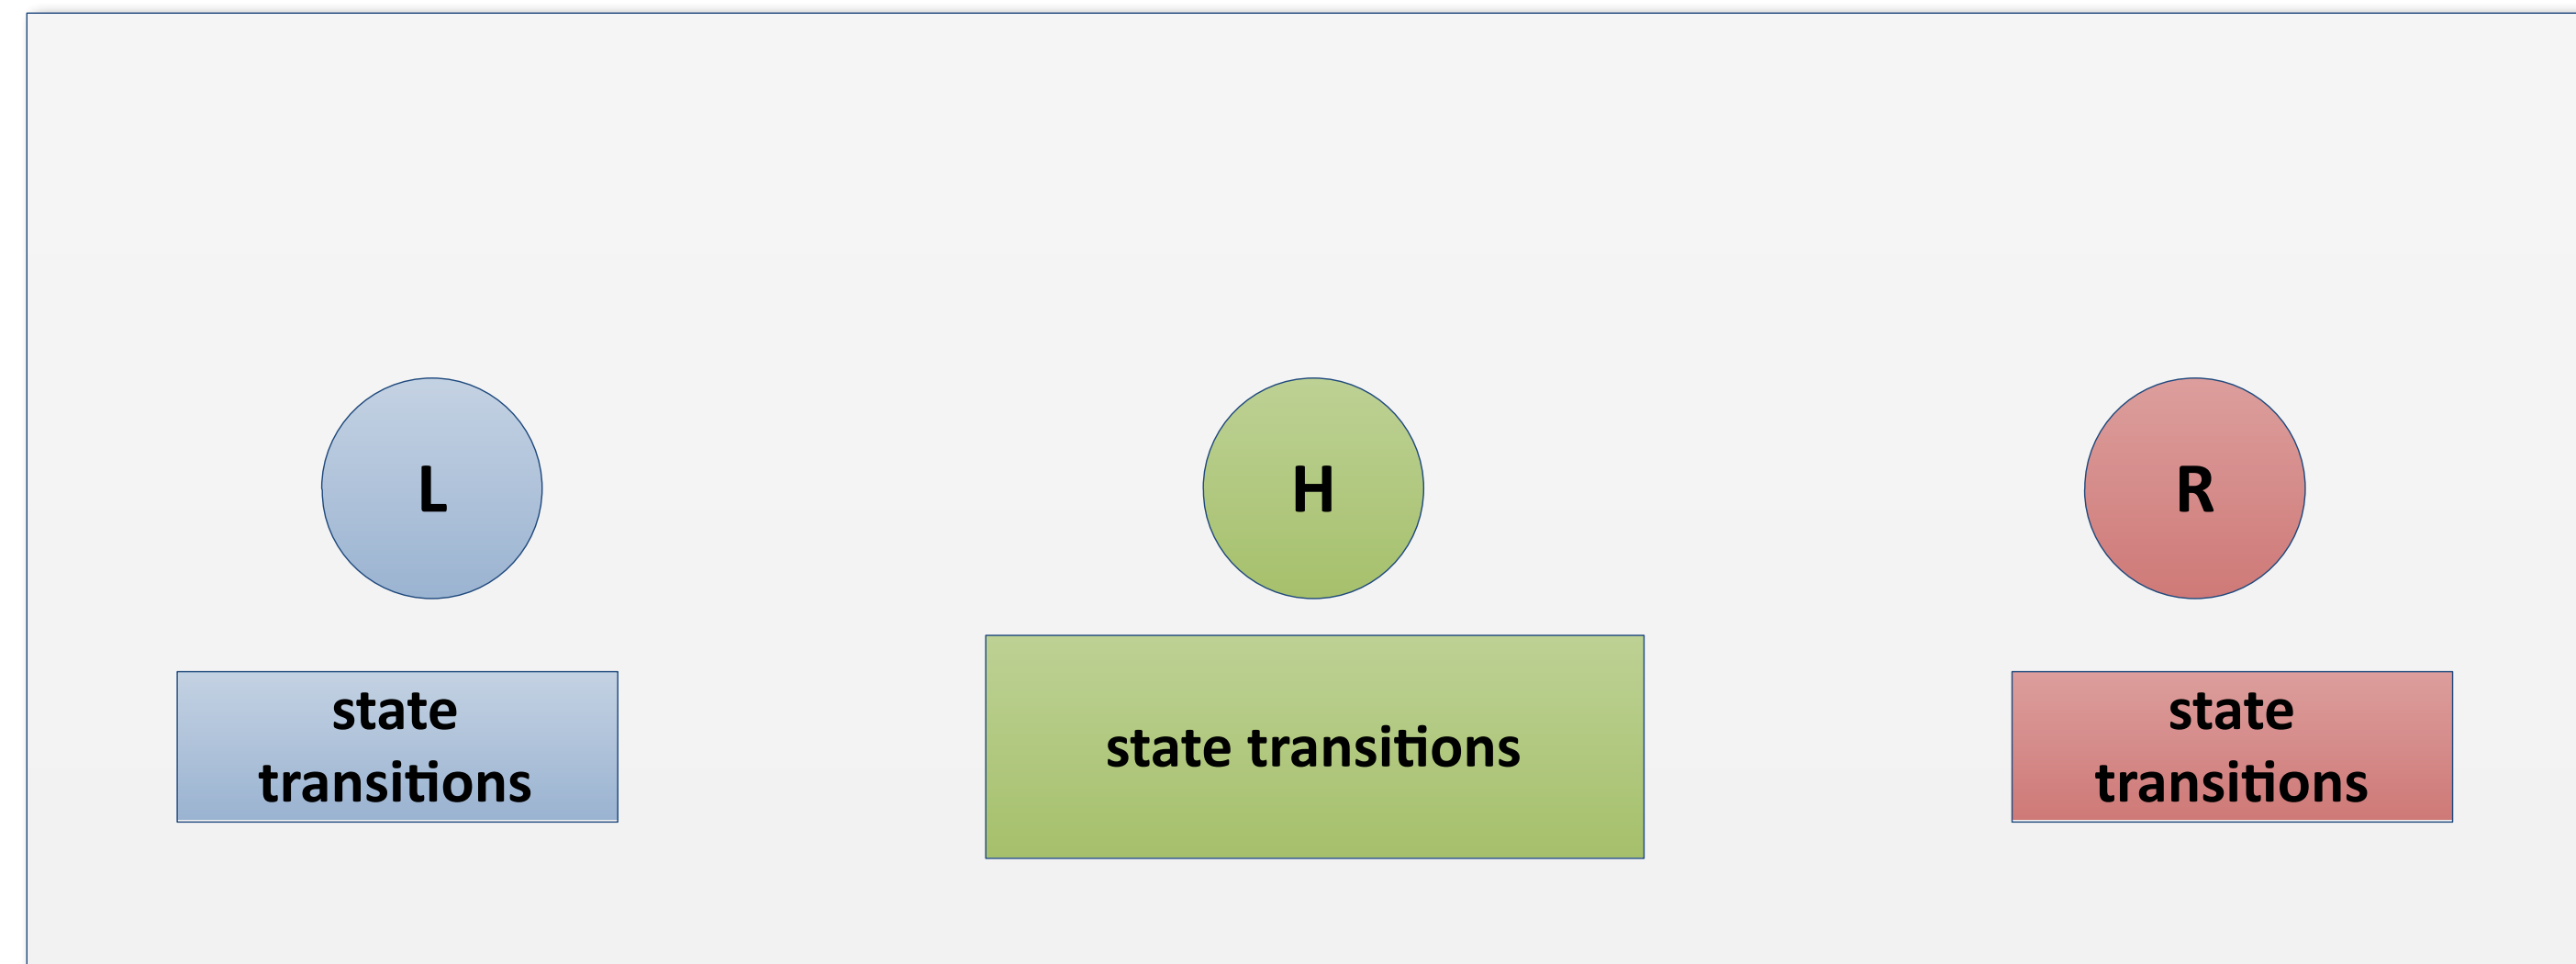
H: Home Node

Node hosting the memory block

R: Remote Node

Any other node participating in the transaction

Assume only MSI states



DIRECTORY EXAMPLE: READ TRANSACTION

L: cache miss on a load instruction

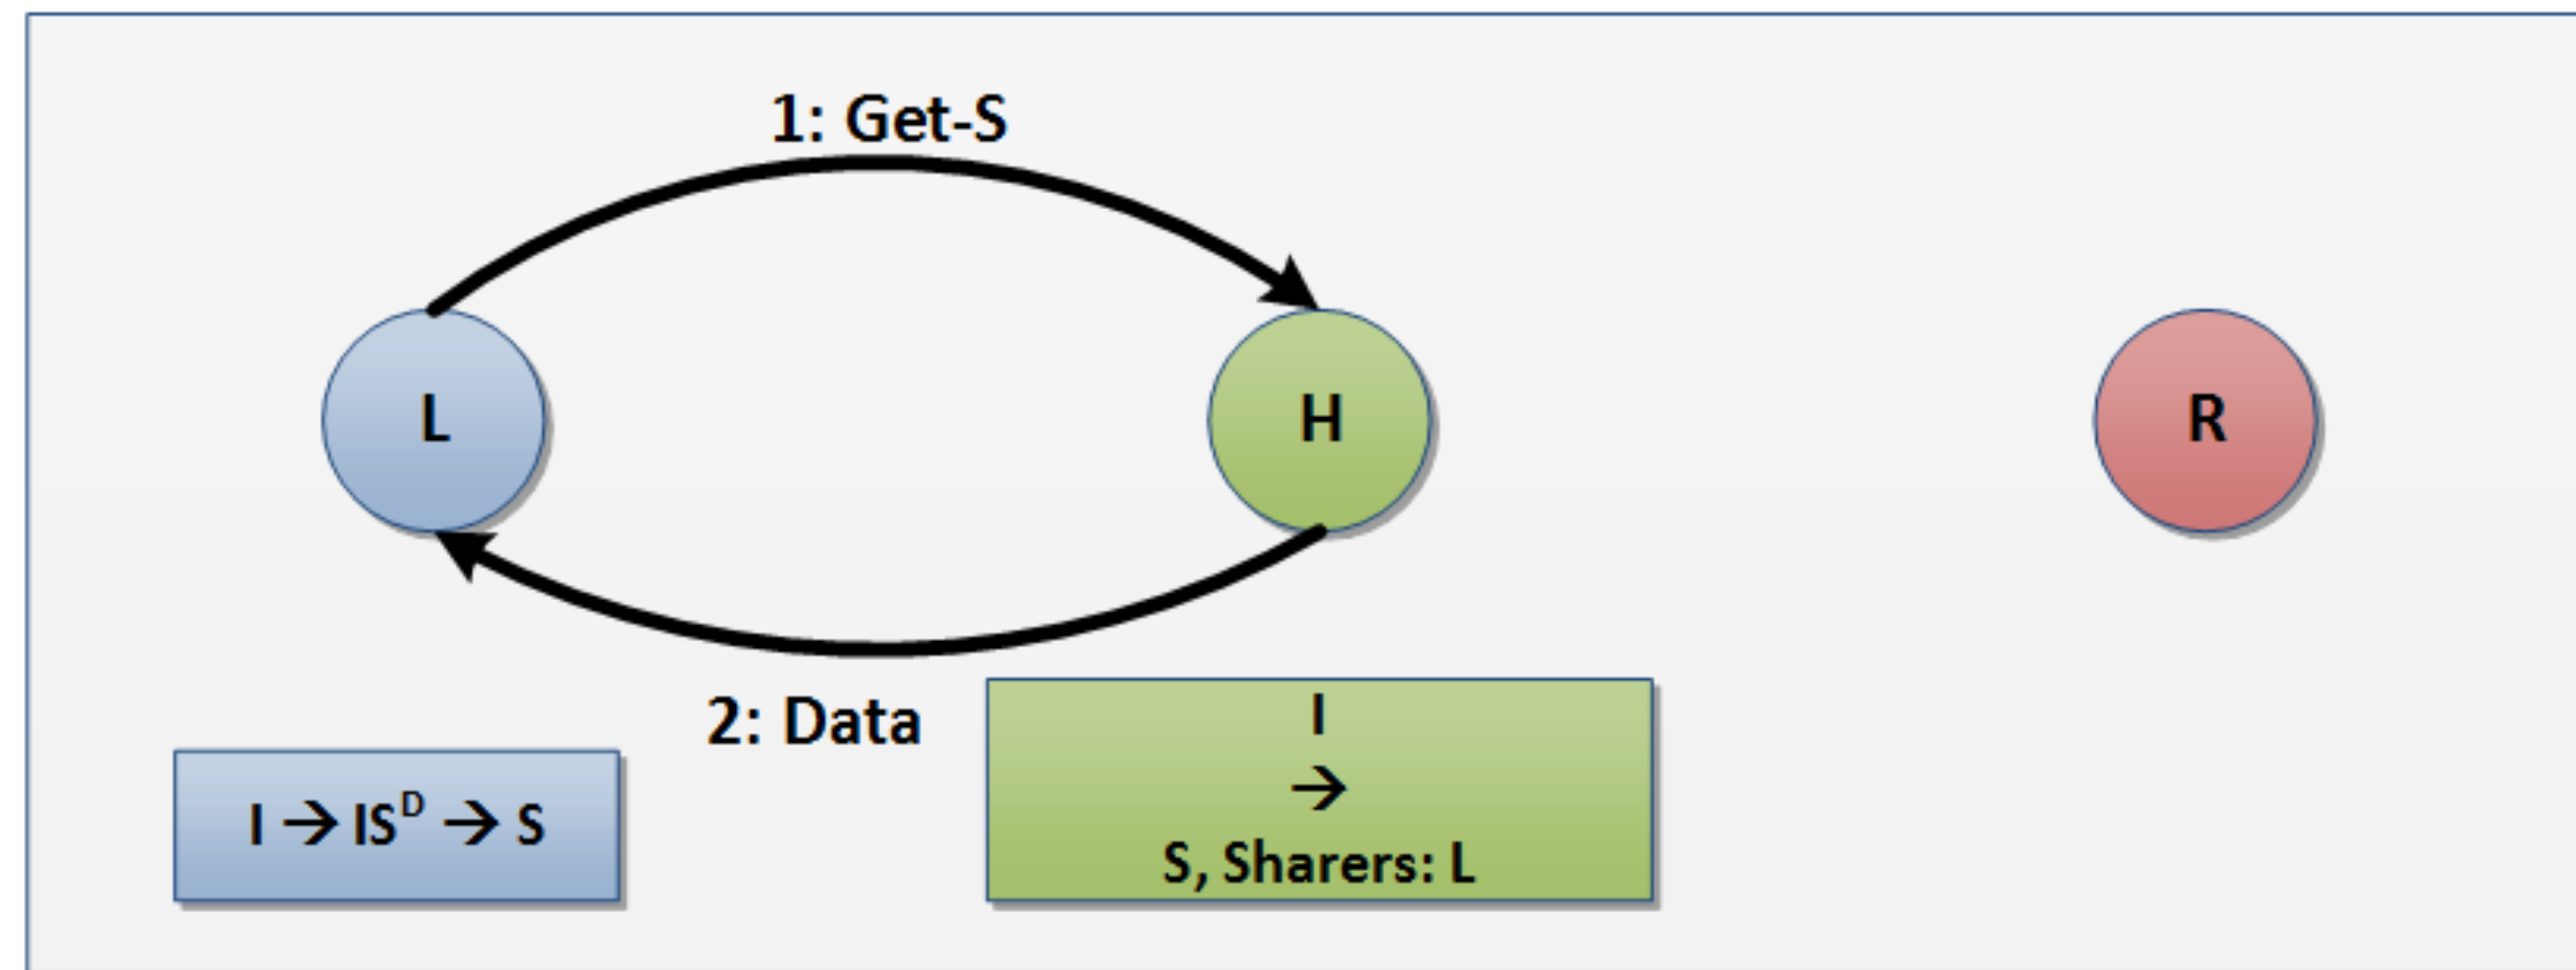
Same if there are other sharers:

Transient states (like IS^D)

H: $\{I, S\} \rightarrow S$

Superscript A: waiting for
acknowledgement

Superscript D: waiting for data

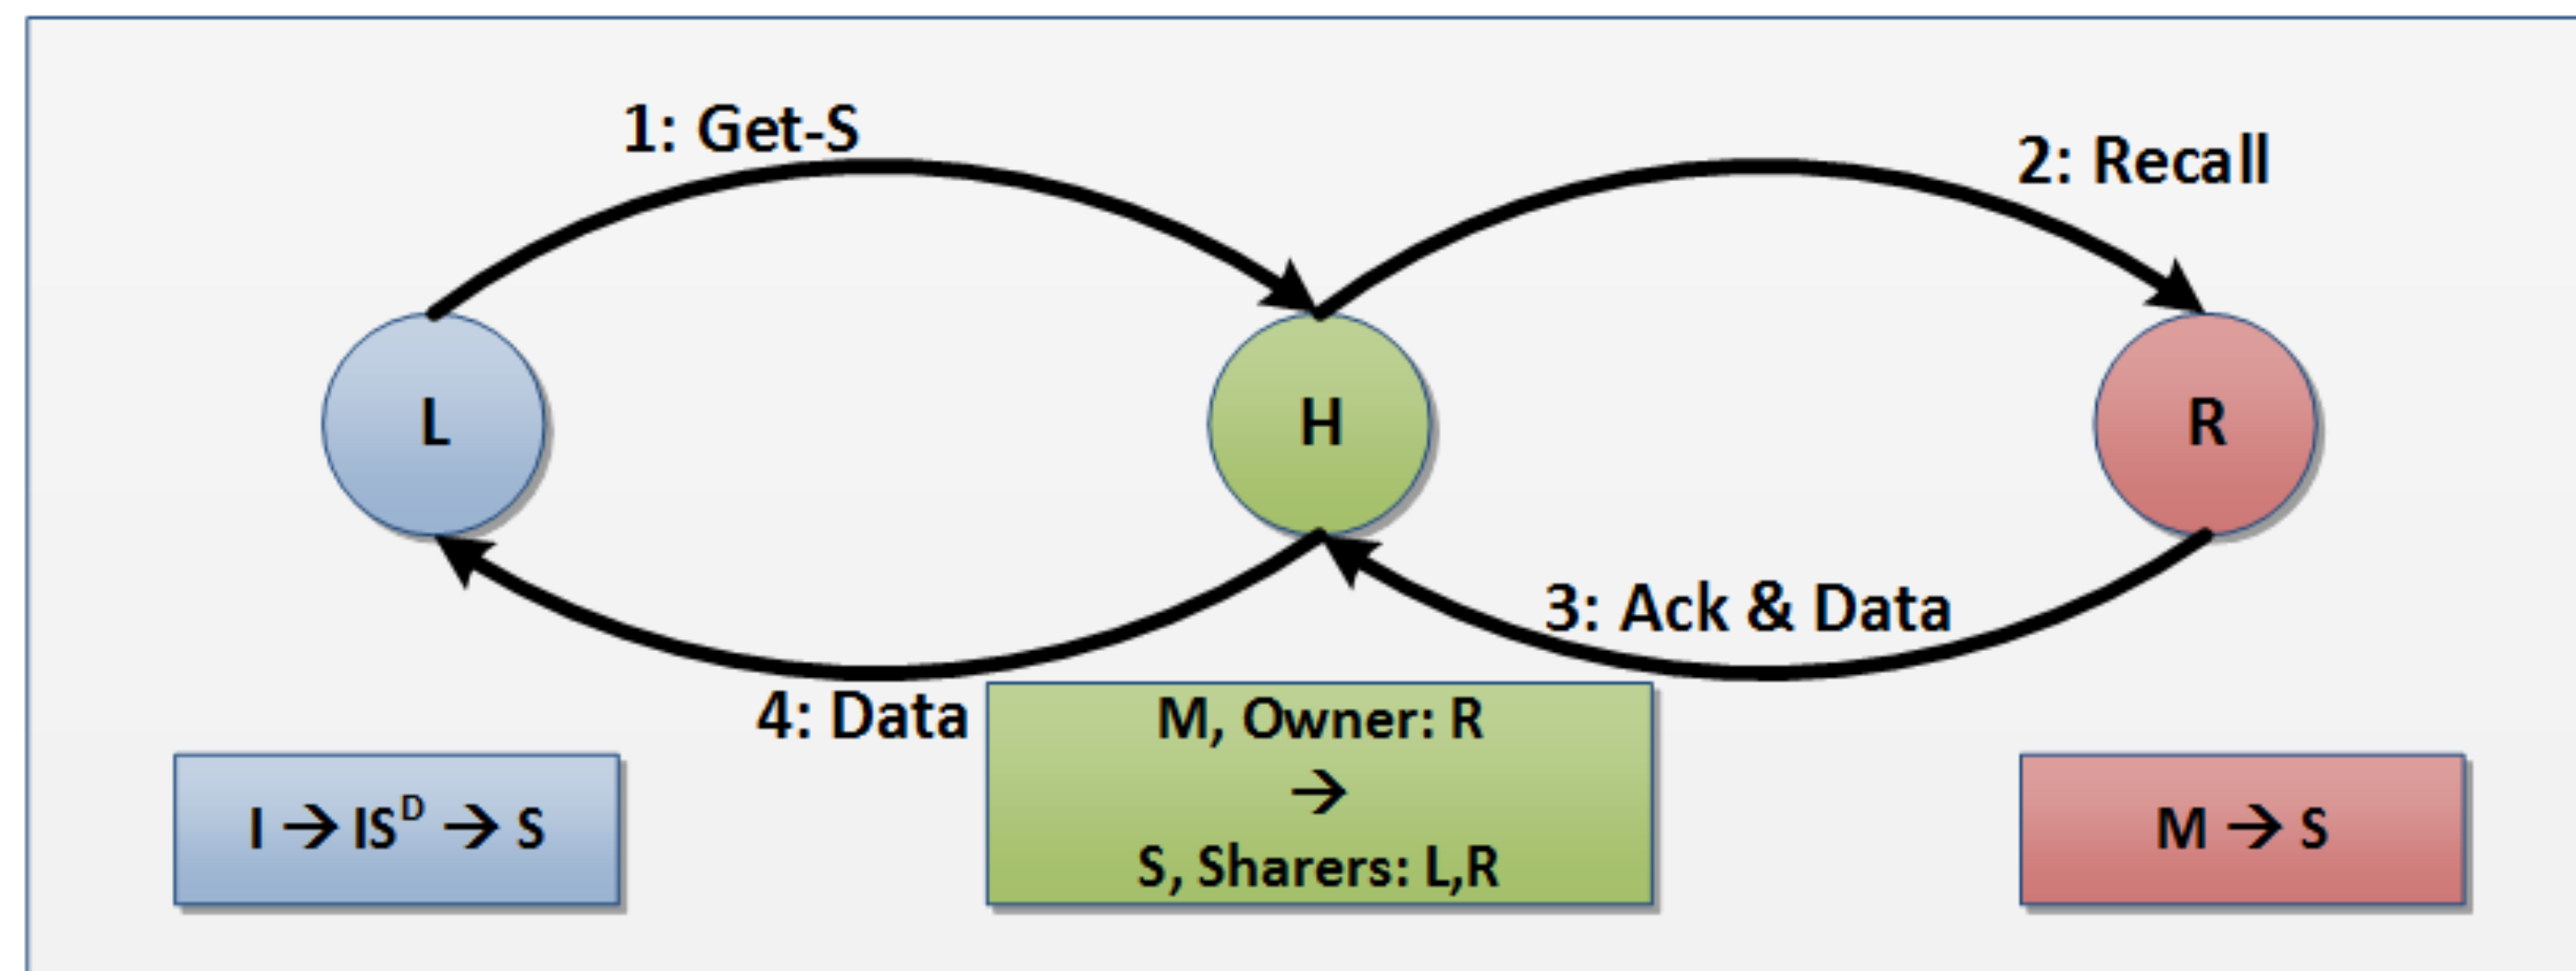


DIRECTORY EXAMPLE: 4-HOP READ TRANSACTION

L: cache miss on a load instruction

Block was previously in modified state at R

Average latency per hop?



DIRECTORY EXAMPLE: 3-HOP READ TRANSACTION

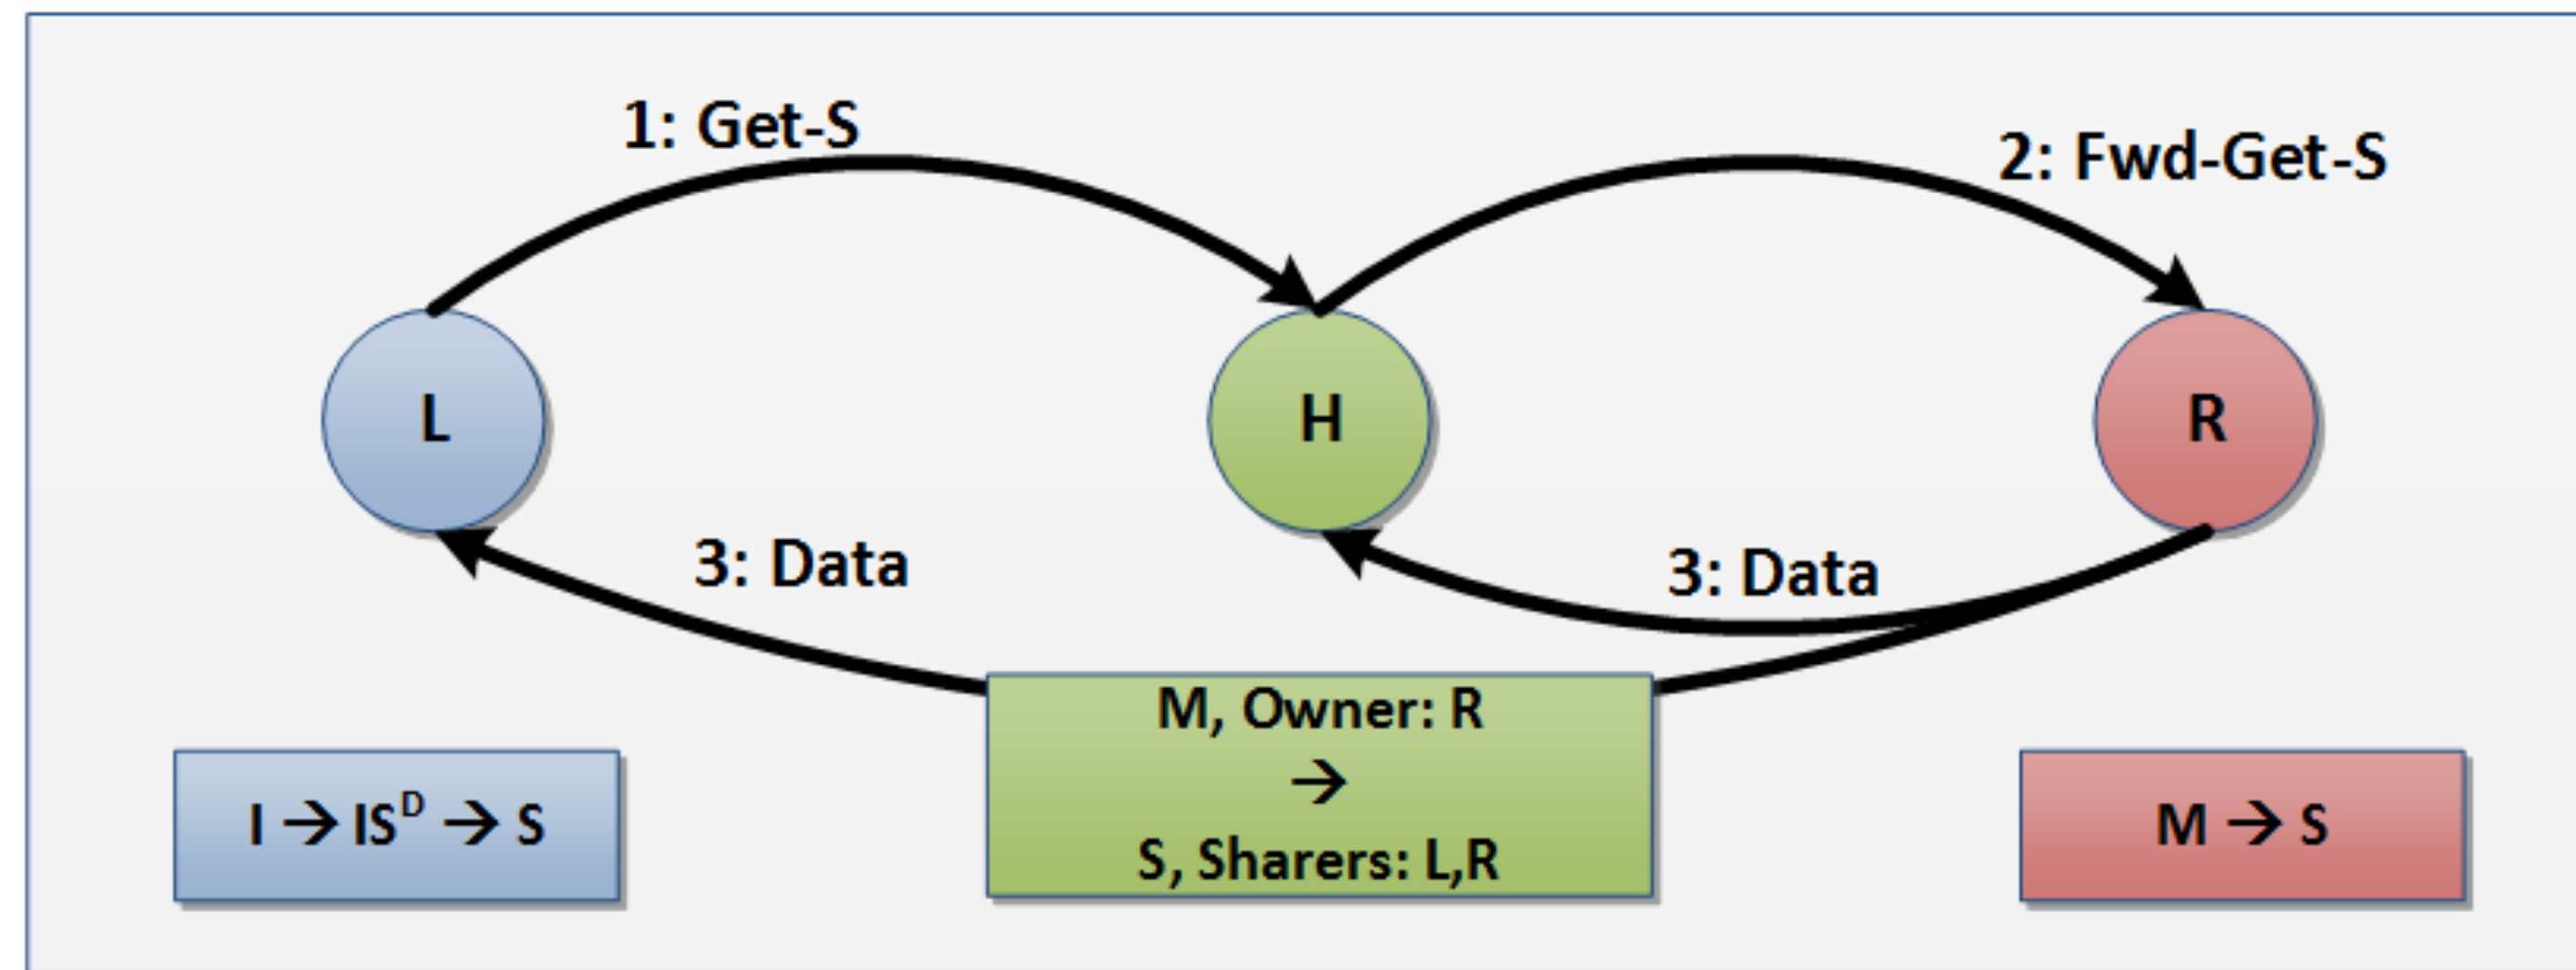
L: cache miss on a load instruction

Block was previously in modified state at R

Optimization: send response directly to L

Why also send data to H?

MSI used here - no O state



DIRECTORY EXAMPLE: 3-HOP WRITE TRANSACTION

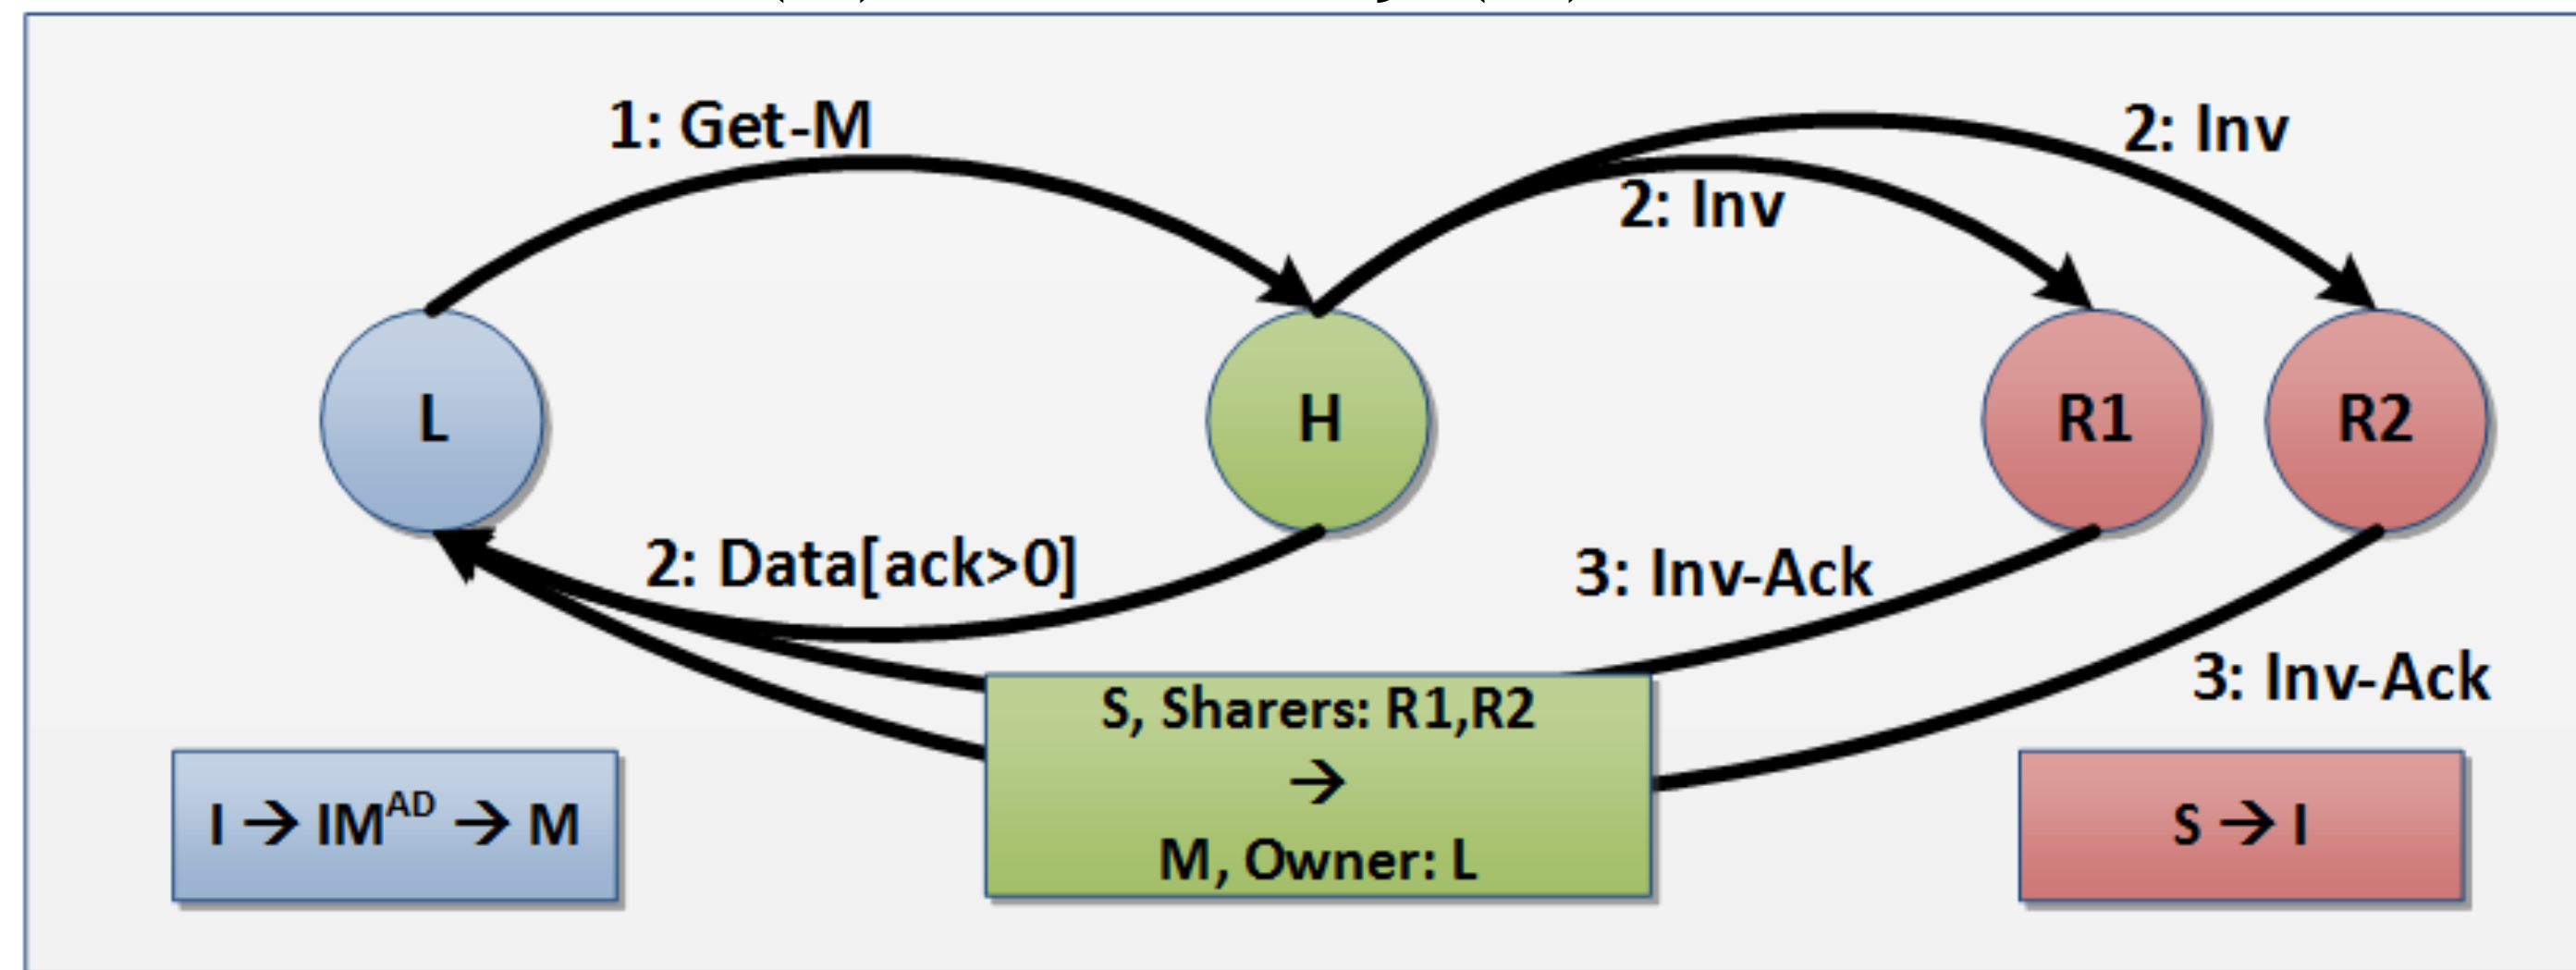
L: cache miss on a store instruction

Block was previously in shared state at R1, R2

Similar applies to L: S \rightarrow M

Note on $\text{ack} > 0$

Options: data from cache (R) or memory (H)



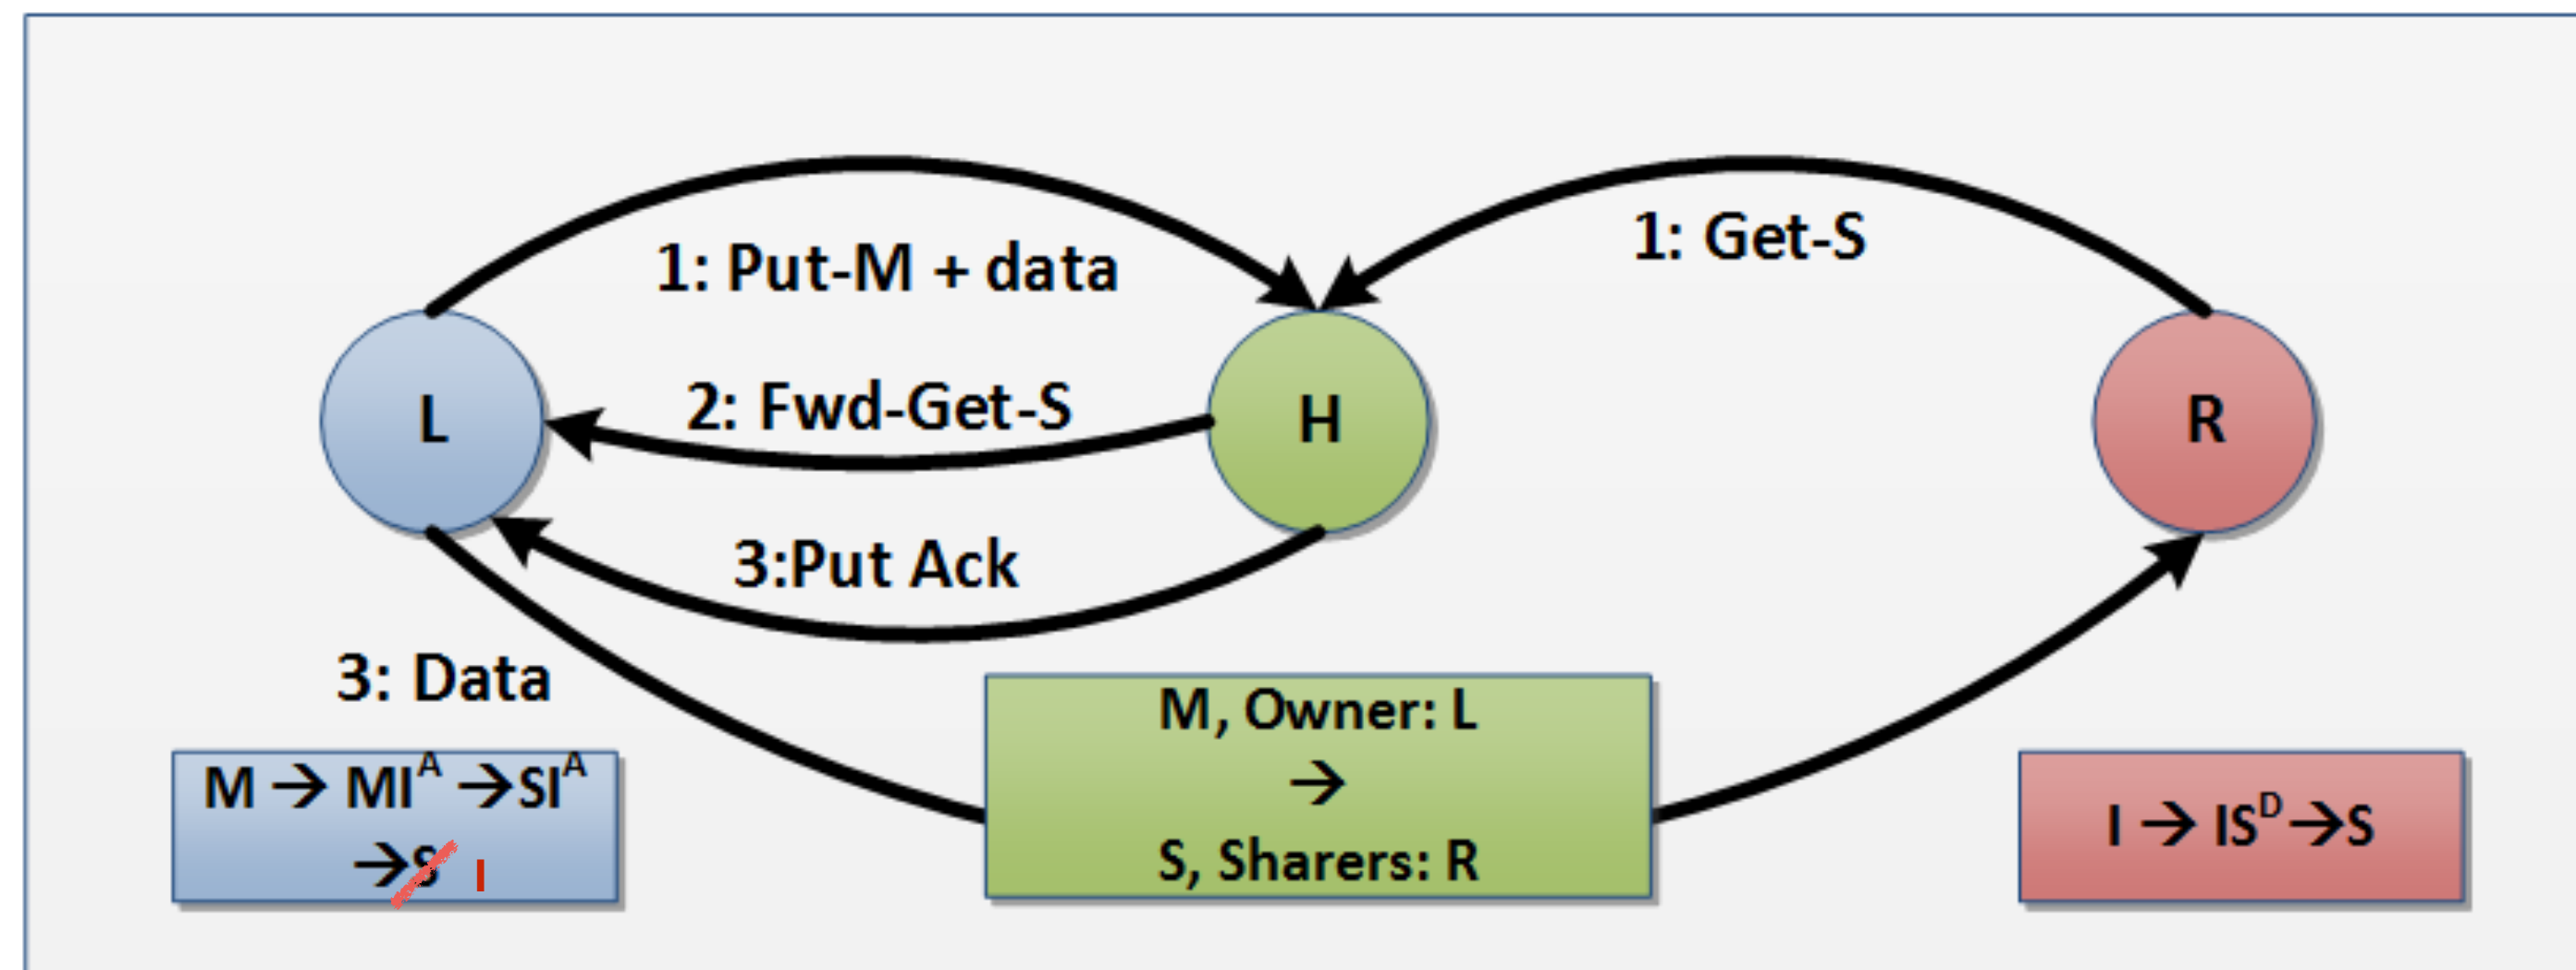
DIRECTORY EXAMPLE: WRITEBACK-READ RACE

L: has dirty copy, wants to write back to H

R concurrently sends a Get-S to H

H rather forwards the Get-S request than serving it, as caches are much faster than a memory access

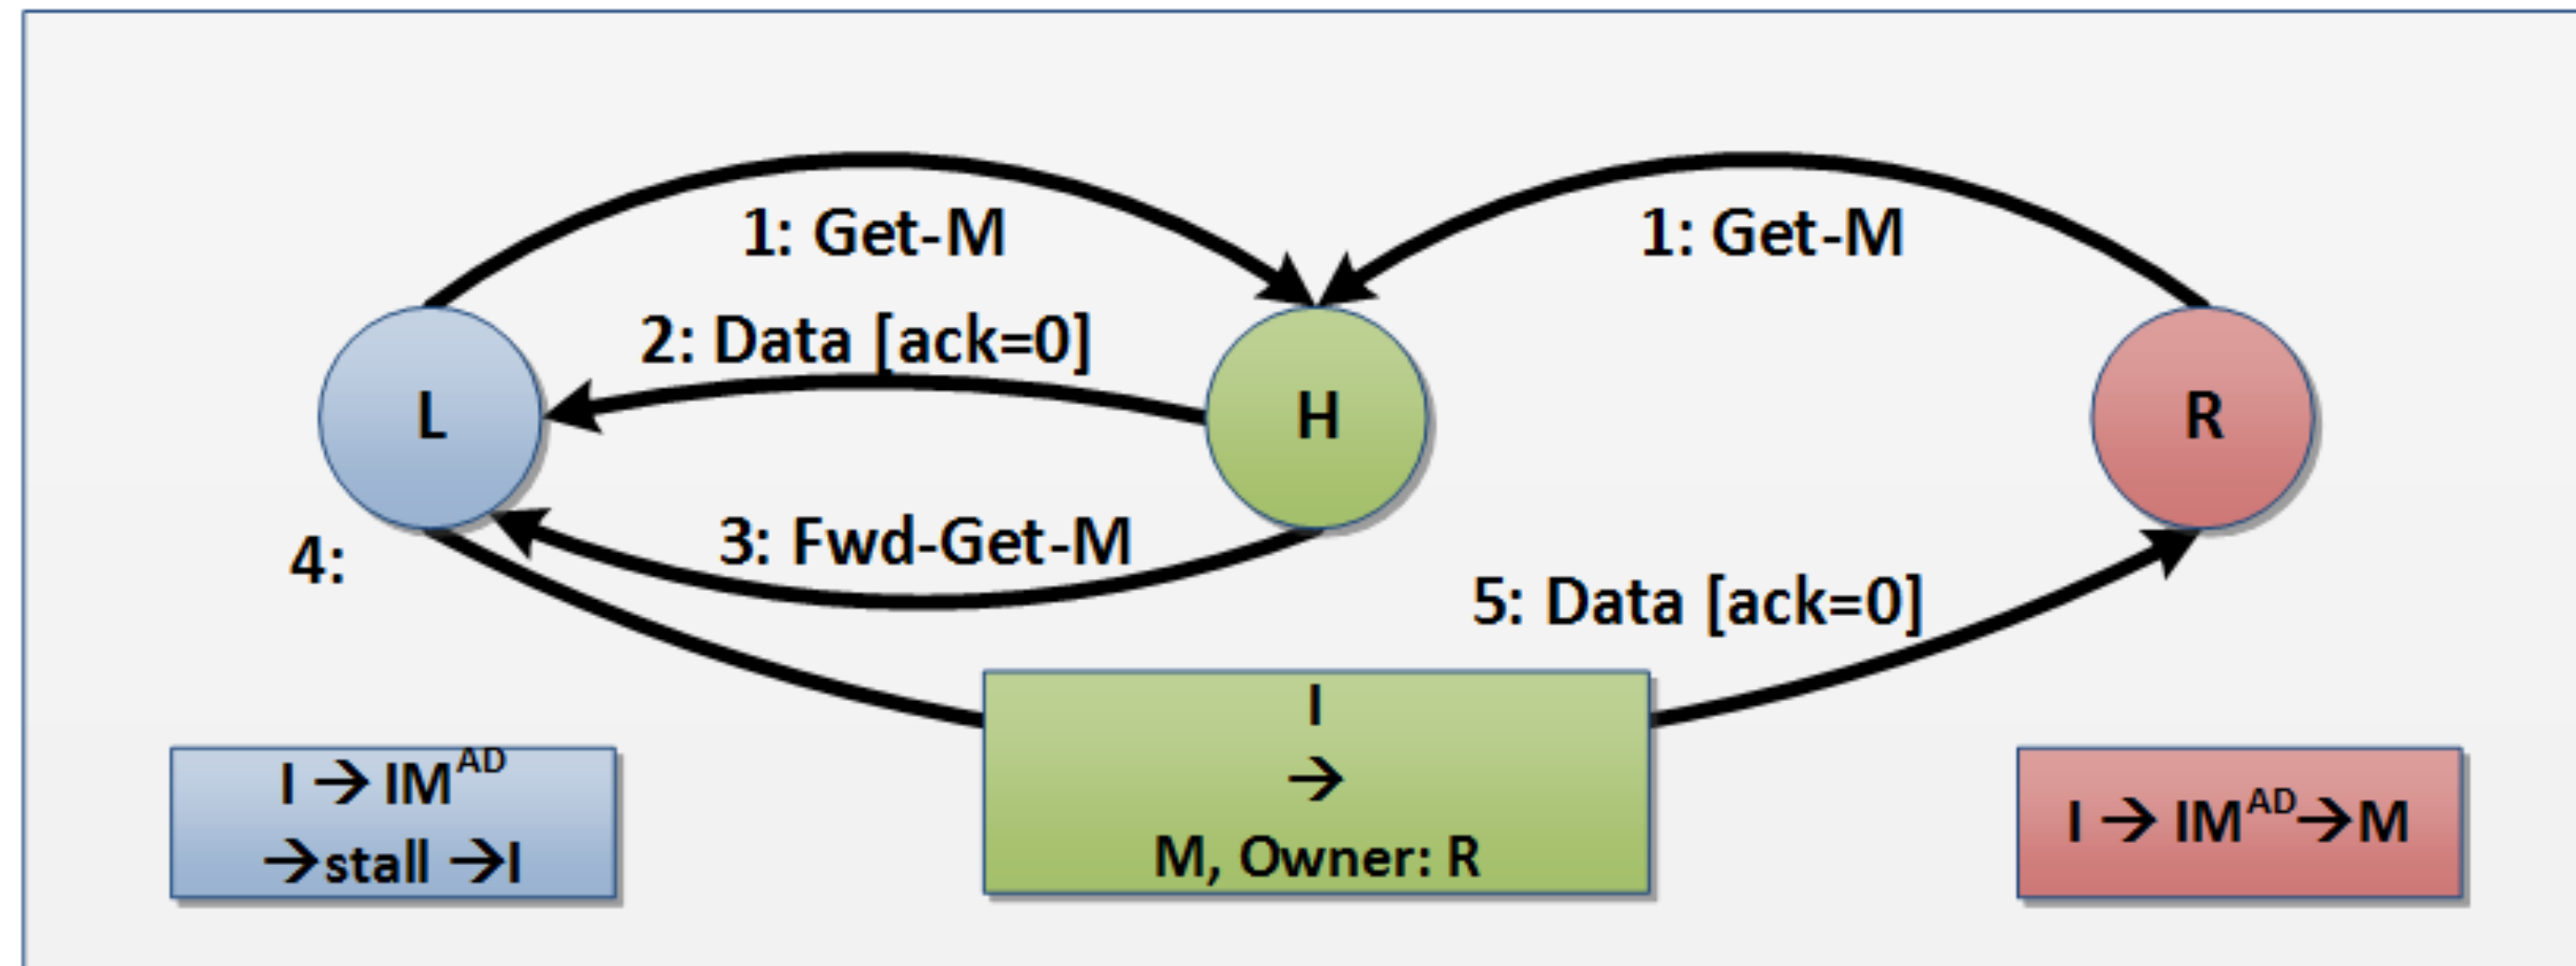
For R sending a write to H: Fwd-Get-M & L: $M \rightarrow MI^A \rightarrow II^A \rightarrow I$



DIRECTORY EXAMPLE: STORE-STORE RACE

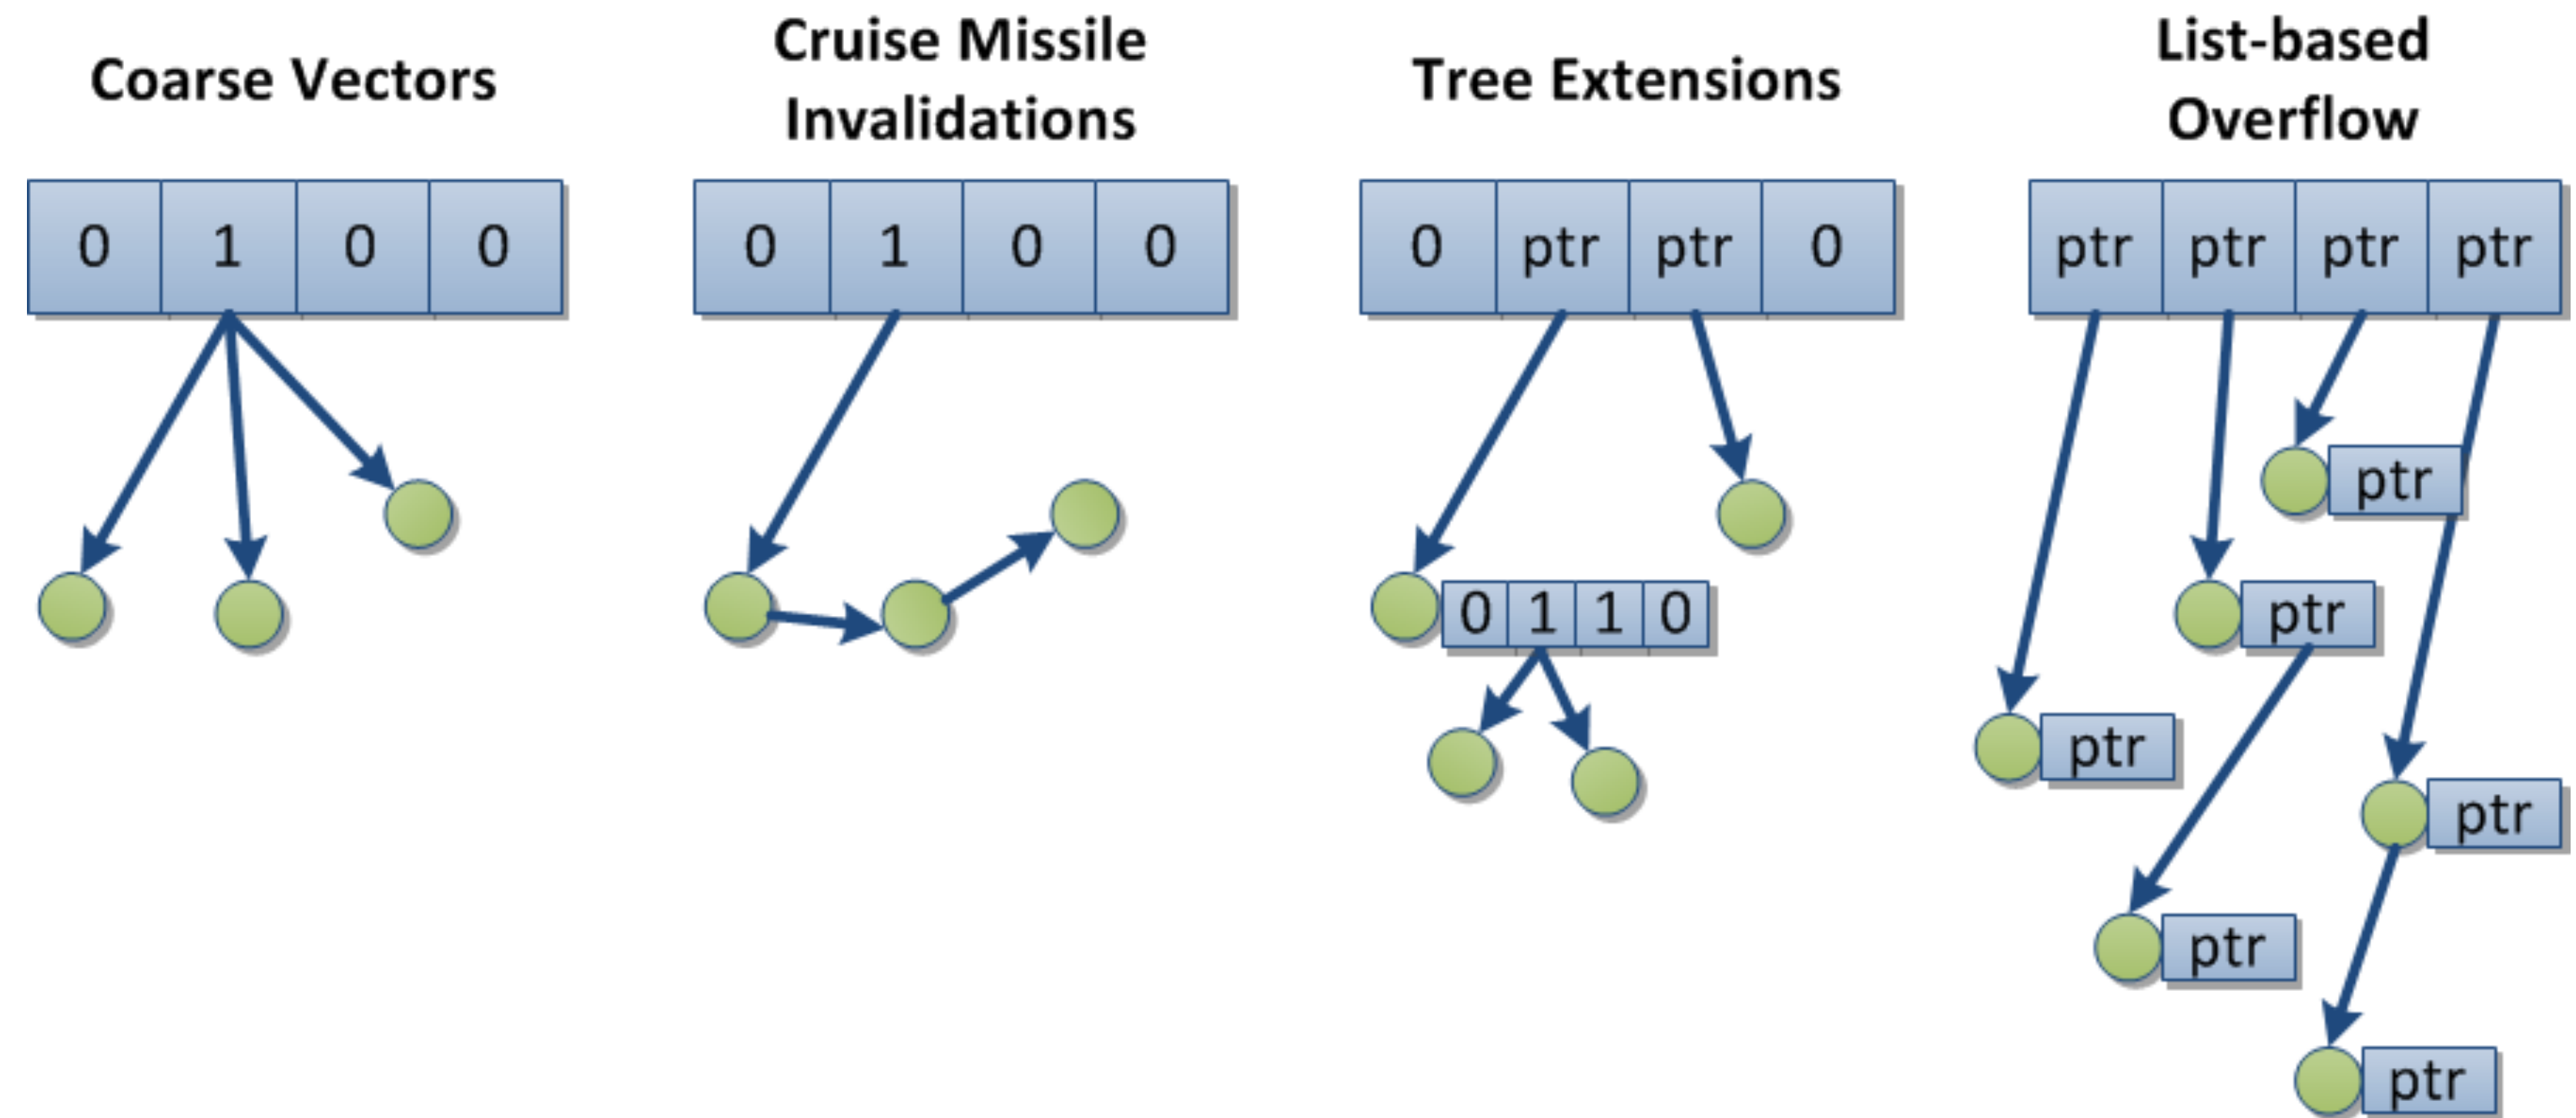
Line is invalid, both L and R race to obtain write permission

L stalls for a single store, then forwards ownership



OPTIMIZATIONS

CONTENTS OF A DIRECTORY: SHARER LIST OPTIMIZATIONS



Example: How to send around an invalidate request

Open question is the number of responses to expect (serialization)

ADVANCED: NOTIFICATION OF CLEAN EVICTION

Should directory learn when clean blocks are evicted? (Put-S)

Avoid broadcast, free pointer in limited pointer schemes (scarce resource) (yes)

Avoids unnecessary invalidate messages later (yes)

Read-only data never invalidated (extra evict messages) (no)

Notification traffic is unnecessary (no)

Avoid some protocol races (yes)

Race example without Put-S

1. Cache silently evicts a block in S
2. Cache sends Get-S request to re-obtain that evicted block in S
3. Before receiving the response, it receives an invalidation

Part of first or second period?

Be safe: always invalidate block when it arrives (more efficient solutions complicate the protocol)

ALTERNATIVE: SPARSE DIRECTORIES

Most of memory is invalid, why waste directory space?

No need to reserve a directory entry for each memory block

Instead, use a directory cache

Only a cached memory block gets an entry

Any address without an entry is invalid

If full, need to evict and invalidate a victim entry

Generally needs to be highly associative

Cache invalidation patterns

Assumption: on a write to a shared location, number of caches to invalidate is typically small

Otherwise: broadcasts would be better than directories

Experience tends to validate this

=> Sparse directory could maintain only a pointer to one node that is caching a block.
Caching block contains pointer to next one.

COMMON SHARING PATTERNS

Code and read-only objects

No problem

Migratory objects

Even as number of caches grow, typically only 1-2 invalidation

Mostly-read objects

Invalidations are expensive but infrequent => ok

Frequently read/written objects (task queues)

Frequent invalidations, hence sharer list is typically short

Synchronization objects

Low-contention locks: few invalidations

High-contention locks: might need special support (e.g., MCS)

Badly-behaved objects

DESIGN PRINCIPLES

DESIGN PRINCIPLES

Think of sending and receiving messages as separate events

At each step, consider what new requests can occur

E.g.: Can a new writeback overtake an older one?

Two messages traversing same direction implies a race

Need to consider both delivery orders

Usually results in a branch in the coherence FSMs

Need to make sure messages can't stick around lost

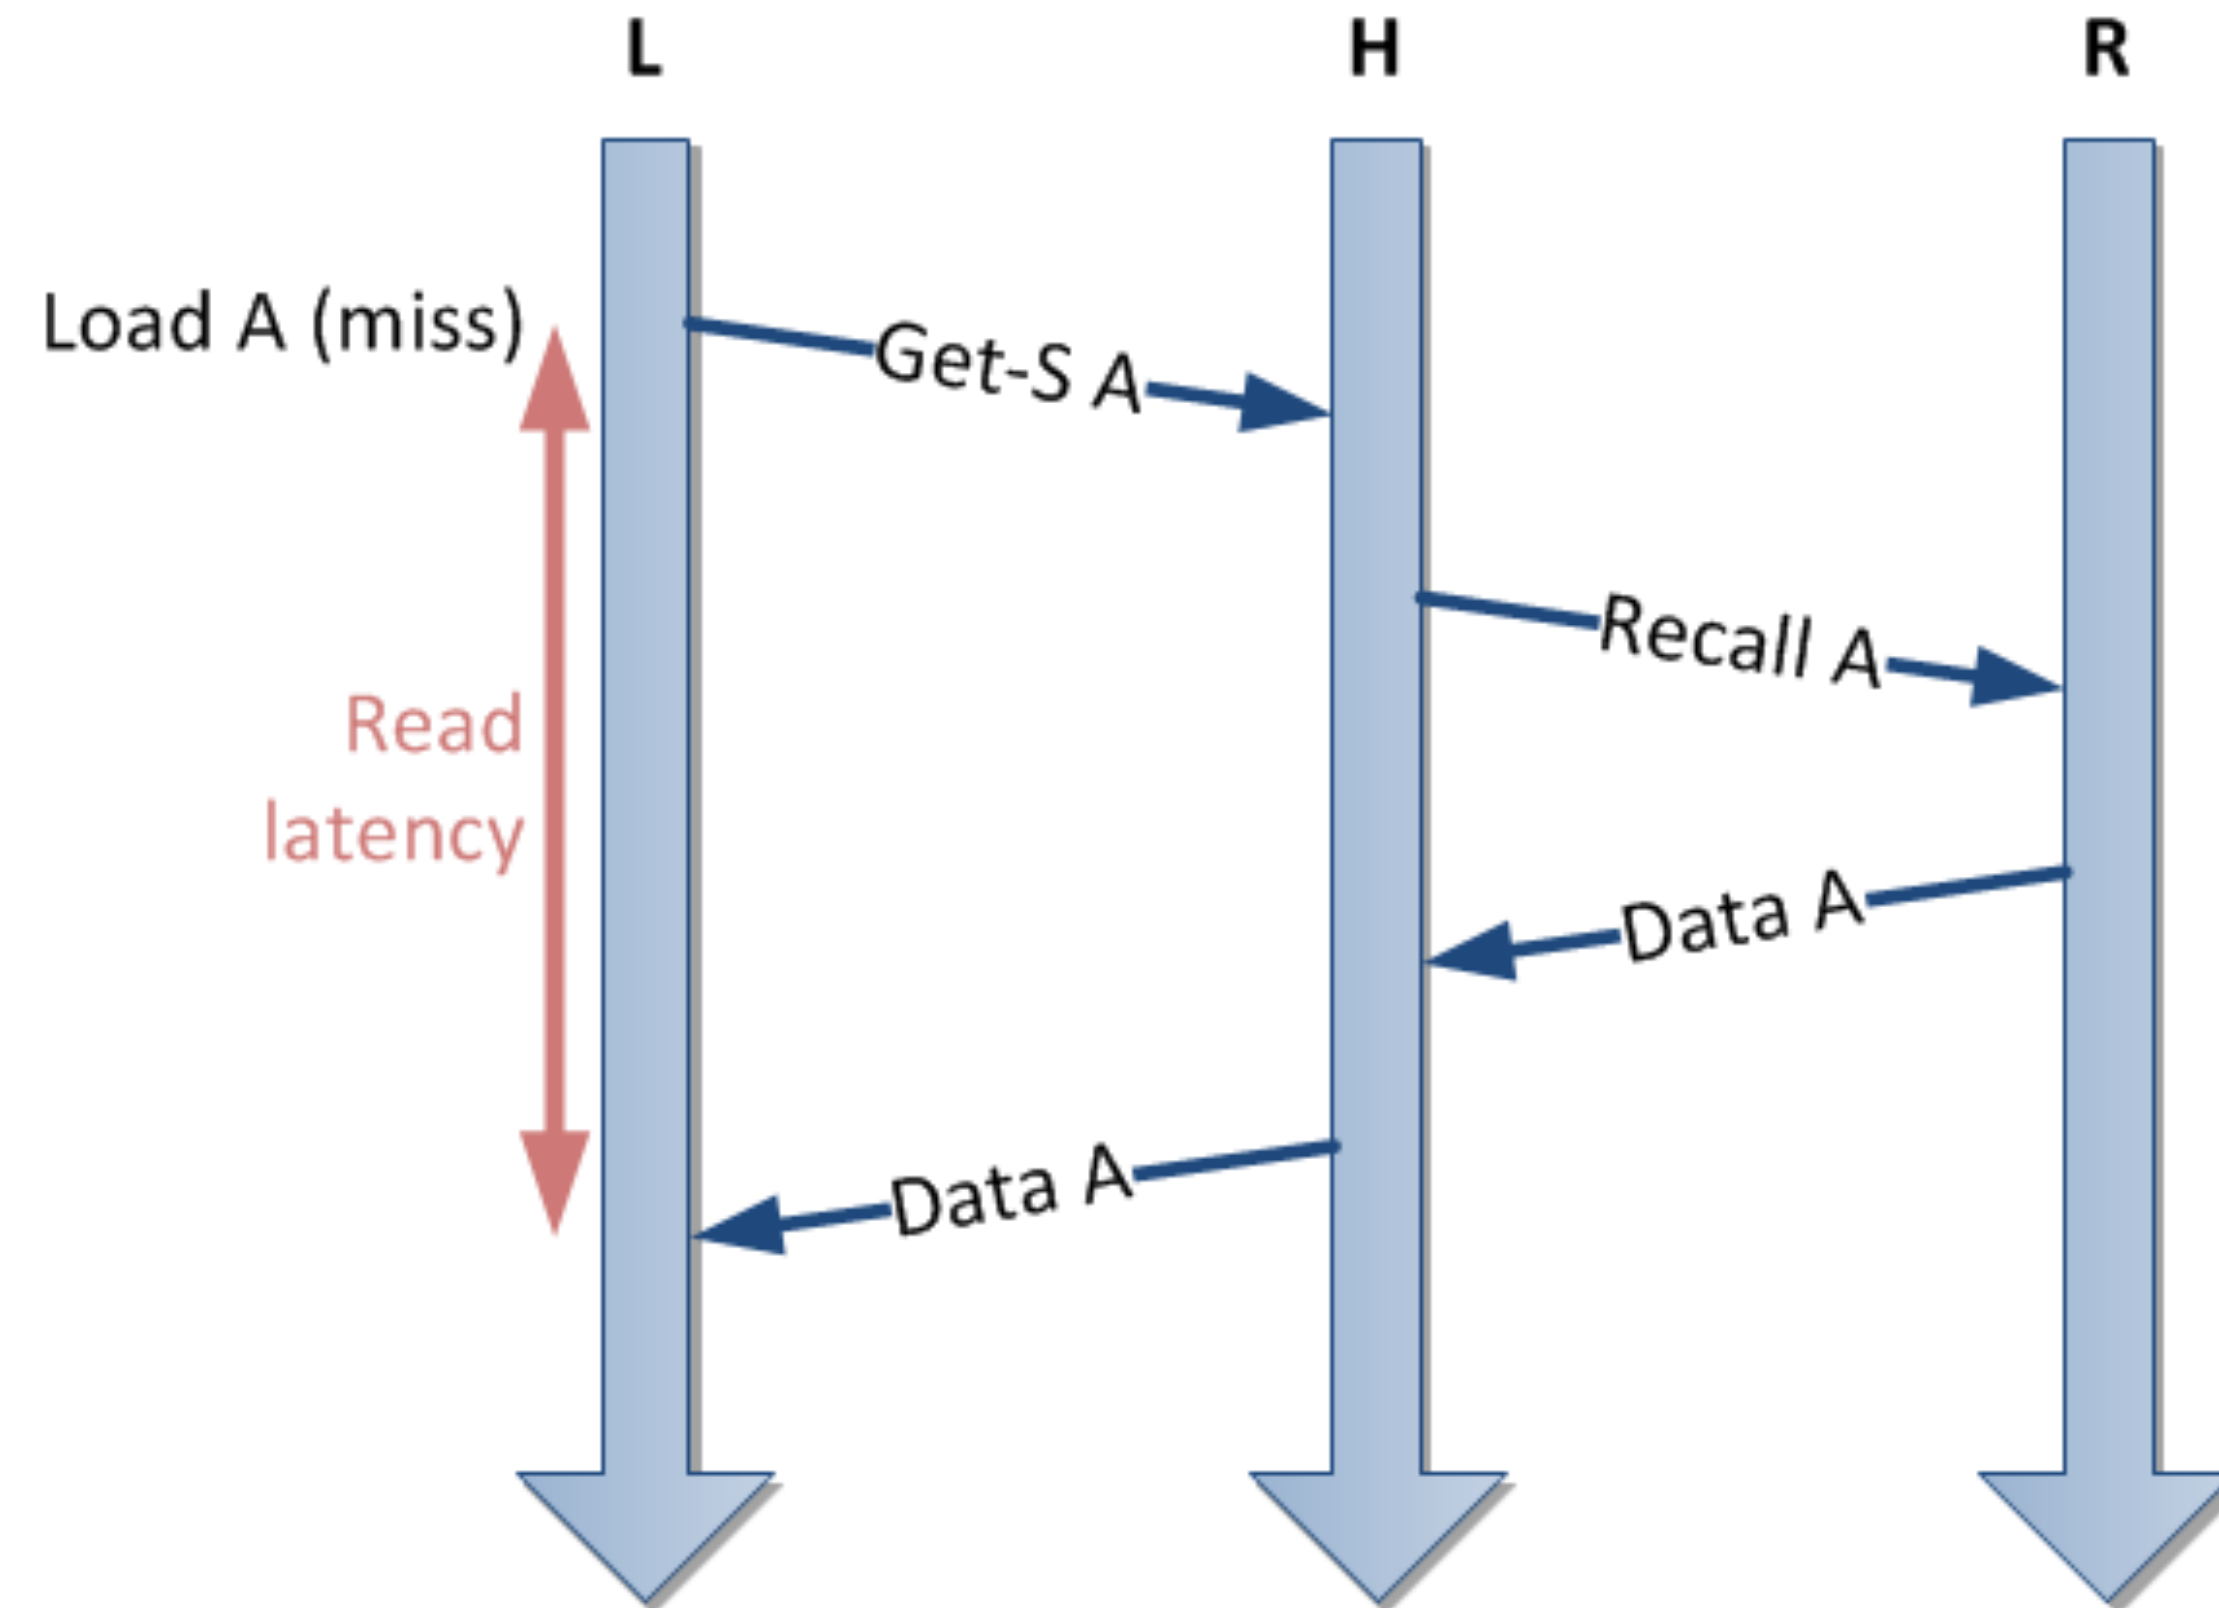
Every request needs an ack

Extra states to clean up messages

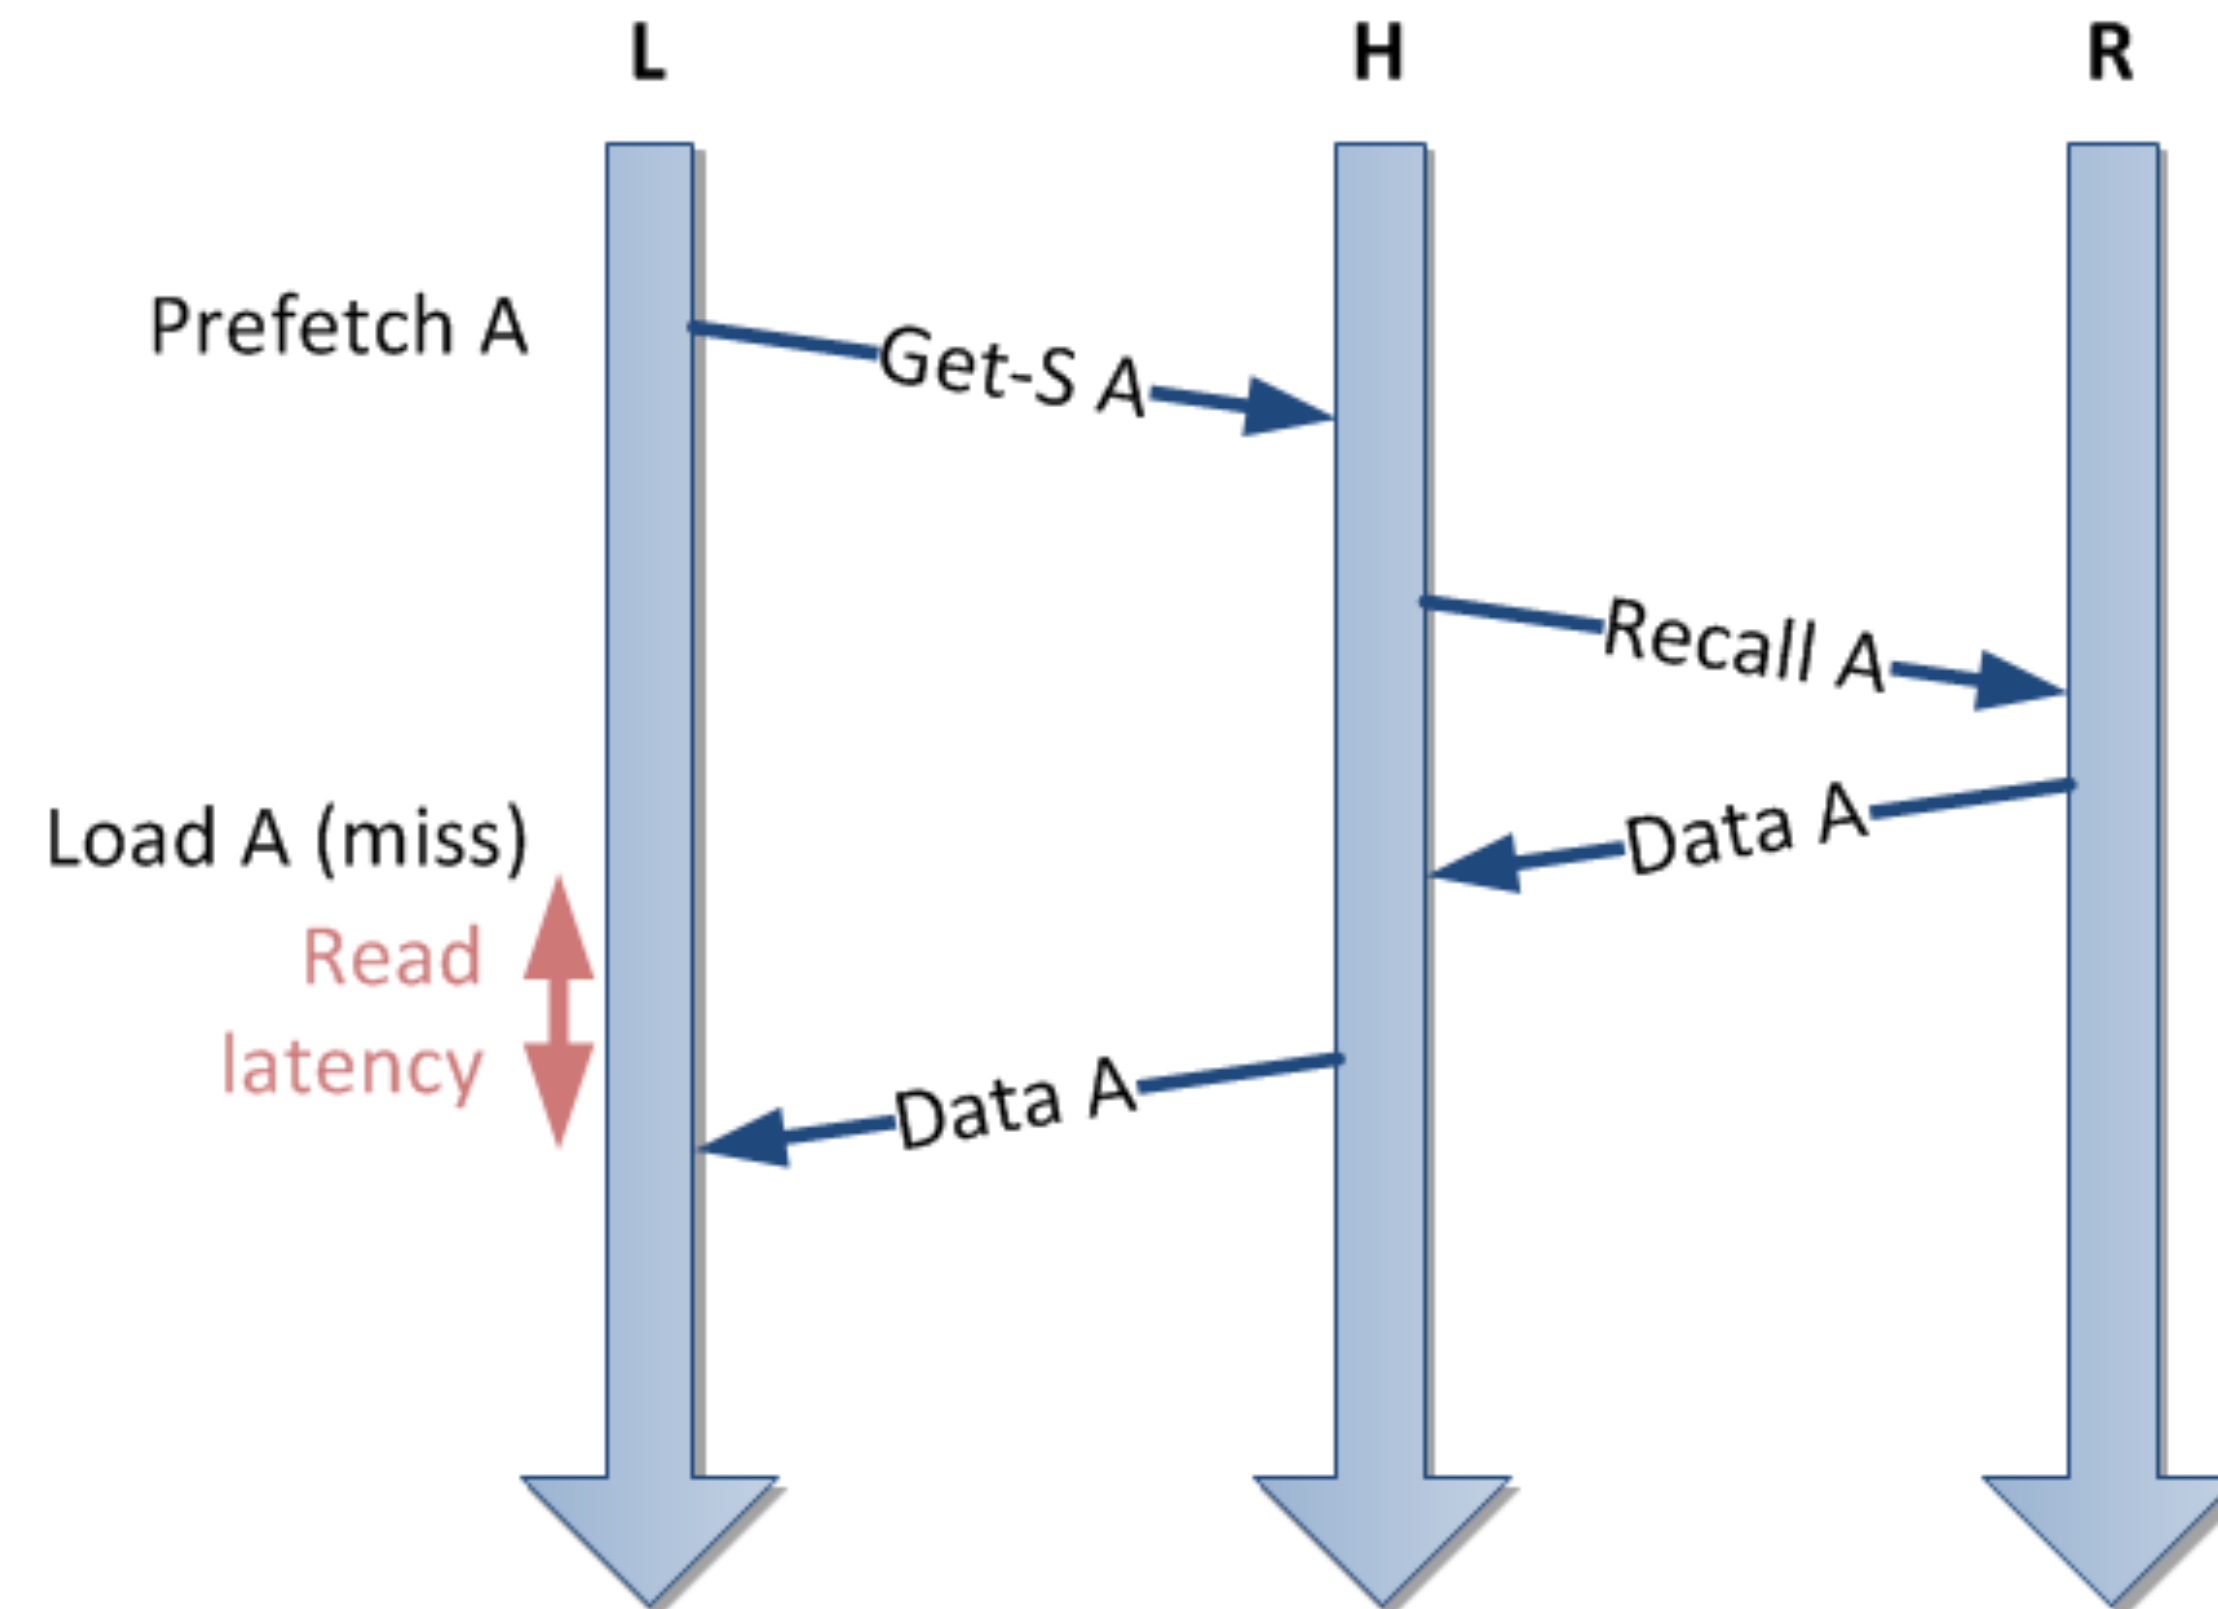
Often, only one node knows how a race resolves

Might need to send messages to tell others what to do

OPTIMIZING COHERENCE PROTOCOLS



OPTIMIZING COHERENCE PROTOCOLS: PREFETCHING



OPTIMIZING COHERENCE PROTOCOLS: MIGRATORY SHARING

Each read/write pair results in read miss & upgrade miss

Coherence FSM can detect this pattern

 Detect via back-to-back read-upgrade sequences

 Transition to “migratory M” state

 Upon a read, invalidate current copy, pass in “migratory E” state

Node 1	Node 2	Node 3
Read X		
Write X		
		Read X
		Write X
	Read X	
	Write X	

OPTIMIZING COHERENCE PROTOCOLS: PRODUCER/CONSUMER

Upon read miss, downgrade instead of invalidate

Detect because there are 2+ readers between writes

More sophisticated optimizations

Keep track of prior readers

Forward data to all readers upon downgrade

Node 1	Node 2	Node 3
Read X		
Write X		
		Read X
	Read X	
Read X		
Write X		
		Read X
	Read X	

OPTIMIZING COHERENCE PROTOCOLS: SHORTCOMINGS

Optimizations have to be included in coherence FSM

- Complex! Adds transitions, states, races...

- Hard to verify even for basic protocols

- Can target only simple sharing patterns

- Can learn only one pattern per address at a time

Each optimization contributes to state explosion

SUMMARY

No free lunch: scalability comes at a certain price

Latency of paramount importance

- Large-scale NUMA systems already suffer a lot from latency issues

Various optimizations possible, but often negative impact on latency due to complexity

Number of sharers is typically small => limited multicast & limited number of directory entries

Main challenges


- Reducing the overhead of directory storage

- Outstanding messages and latency

However/thus, many machines today employ directories

- Examples: AMD's probe filter (HT), Intel's snoop filter (QPI), SGI Altix, NUMASCALE (successor of SCI), ...

EXAMPLE: INTEL CORE I7 PROCESSOR



A ring is not a bus

Centralized directory for all blocks in L3

Inclusion property is important

Directory maintains list of L2 caches containing block

=> Multicast of coherence messages to L2s that contain block

Instead of broadcast

APPENDIX: PROTOCOL ISSUES

CACHE COHERENCE PROTOCOL ISSUES (1)

NACKs are used to re-issue requests

E.g., when racing

Does the protocol use negative acknowledgements (retries)?

NACKs are a bad idea for CC protocols (livelock, starvation, fairness)

Usually a consequence of protocol interaction



CACHE COHERENCE PROTOCOL ISSUES (2)

Is the number of active messages (sent but unprocessed) for one transaction bound?

Scalability issue - how much queue space is needed

MSHR - Miss-Status Handling Registers

Example: coarse-vector vs. cruise-missile invalidation



CACHE COHERENCE PROTOCOL ISSUES (3)

Does the protocol require clean eviction notifications (Put-S)?

(Silent evict) -> Get-M -> Inv ?

How/when is the directory accessed during transactions?

Directory in main memory: latency issues

How to deal with transient states?

Keep it in the directory: unlimited concurrency

Keep it in a pending transaction buffer (e.g., transaction state register file):
faster, but limited pending transactions

Occupancy free: upon receiving an unsolicited request, can directory determine final state solely from current state?

CACHE COHERENCE PROTOCOL ISSUES (4)

How many lanes are needed to avoid deadlocks?

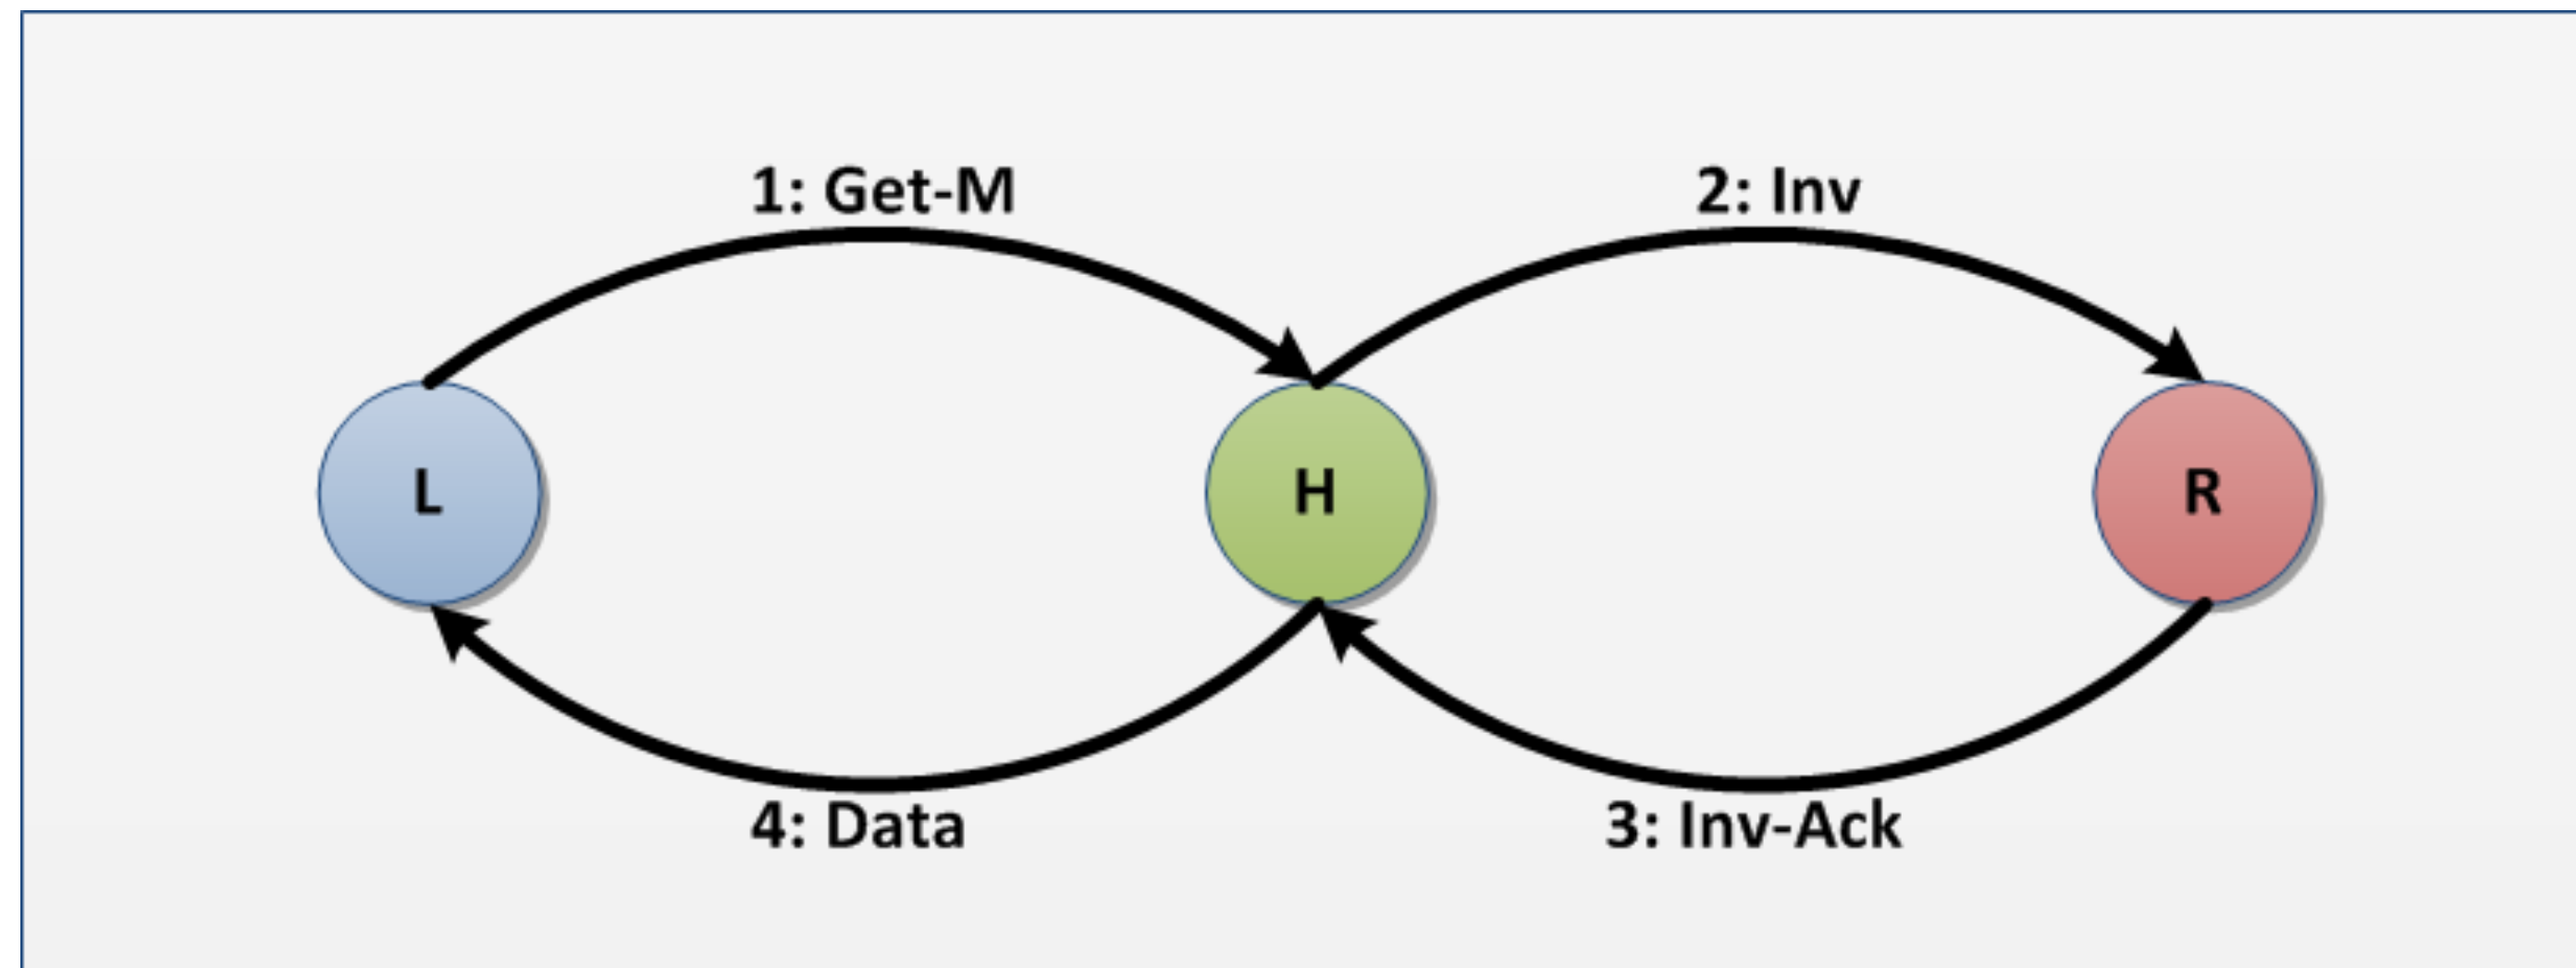
At least two to avoid protocol deadlocks

How to assign lanes to message types?

Secondary (forced) requests may not be blocked by new requests

Replies (completing a pending transaction) may not be blocked by new requests

More may be needed by IO, complex forwarding



CACHE COHERENCE PROTOCOL ISSUES (5)

All messages should be acknowledged

Requests elicit replies

Maximum number of potential concurrent messages for one transaction should be small and constant

I.e., independent of number of nodes in the system

Minimize traffic, minimal interaction, low talk/silence ratio

Use context information to avoid NACKs