

ADVANCED PARALLEL COMPUTING

LECTURE 09 - MEMORY CONSISTENCY

Holger Fröning
holger.froening@ziti.uni-heidelberg.de
Institute of Computer Engineering
Ruprecht-Karls University of Heidelberg

FINAL COHERENCE NOTES

Main coherence problem: overhead and scalability

- Bandwidth for probing / latency increase

- No caches, no coherence problem!

Intuition says load should return latest value

- What is latest?

View from above the clouds (only):

- Coherence is about reads, consistency is about writes

- Coherence concerns only one memory location

- Consistency concerns apparent ordering for all locations

RELAXED MEMORY CONSISTENCY


Consistency model defines correct behavior

- Contract between system and programmer (similar to ISA)

- Restricts ordering of loads/stores

Software-visible

- Coherence is not visible to software



<http://www.linuxjournal.com>

WHY COHERENCE \neq CONSISTENCY

P0

```
A = 1;  
B = 1;  
flag = 1;
```

P1

```
while (!flag);  
print A;  
print B;
```

Intuition says: $A=B=1$

However, not for all consistency models

Coherence doesn't say anything about that

Why?

What does it guarantee in this example?

WHY ARE MEMORY ACCESSSES RELAXED?

Memory accesses are the slowest operations of a processor
(processor/memory gap)

- Huge caches are mandatory to address this disparity

- Caches are partitioned into (almost) independent banks

- E.g.: odd-numbered addresses into cache bank #0, even-numbered into #1

- Idle/busy banks change the visibility of writes

Huge overhead to maintain ordering of memory accesses

=> Performance degradation for strict consistency

=> Only reason for relaxed consistency: performance

STORE ATOMICITY

Again, caches:

P0	P1	P2
<code>A = 1;</code>	<code>while (!A);</code> <code>B = 1;</code>	<code>while (!B);</code> <code>print A;</code>

Intuition says: P2 prints A=1

Key issue here: store atomicity

Does a new value reach all nodes at the same time?

Many (past) commercial systems allow P2 to print A=0

ONE PROGRAM - VARIOUS ORDERINGS

What the user wants: memory references specified in high-level language

C, Java, etc.

Program order: what the compiler makes out of it

Compilers increasingly reorder memory accesses depending on the dependencies they see (see keyword volatile)

Execution order: how the processor executes these instructions

Modern CPUs are microcode-based, out-of-order, superscalar, multi-core architectures

Perceived order: how the processor(s) observe memory operations

Own and others' memory operations are reordered due to caching, interconnect and memory-system optimizations.

Different processors can even have a different view of memory operations

SOME THINGS YOU CAN COUNT ON

1. One processor always sees its own memory operations in program order

Thus, reordering issues arise only between processors

2. An operation is reordered with a store only for different addresses
3. Aligned simple loads and stores are atomic
4. Any relaxation comes with a safety net
I.e., revert to sequential consistency as needed

STRICT CONSISTENCY MODELS

STRICT CONSISTENCY

Strict (or atomic) consistency:
“A read returns the most recent
value written to a location”

Simplest and strictest consistency
model

Relies on a global clock to
serialize all memory operations

Only used for uniprocessors

P0:	W (x) 1		
P1:		R (x) 1	R (x) 1
Valid			


P0:		W (x) 1	
P1:	R (x) 0		R (x) 1
Valid			

P0:	W (x) 1		
P1:		R (x) 0	R (x) 1
Invalid			

SEQUENTIAL CONSISTENCY (SC)

Sufficient condition for SC:

“A multiprocessor is sequentially consistent if the result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program” - Lamport, 1979



SEQUENTIAL CONSISTENCY

1. Every processor issues memory requests in program order

P0:	W (x) 1		
P1:		R (x) 0	R (x) 1

Valid

2. Memory operations happen (start and end) atomically

P0:	W (x) 1		
P1:		R (x) 1	R (x) 2
P2:		R (x) 1	R (x) 2
P3:	W (x) 2		

Valid

3. Must wait for a store to complete before issuing next operation

4. After a load, issuing processor waits for load to complete, before issuing next operation

P0:	W (x) 1		
P1:		R (x) 1	R (x) 2
P2:		R (x) 2	R (x) 1
P3:	W (x) 2		

Invalid

=> Easily implemented with a shared bus:
synchronization point, serializing all accesses

SEQUENTIAL CONSISTENCY

Example code

P0 :	P1 :	P2 :
$W_0(x, 1)$	$W_1(y, 2)$	$R_2(y, 2)$
	$R_1(a, 0)$	$R_2(x, 0)$
		$R_2(x, 1)$

Perceived order

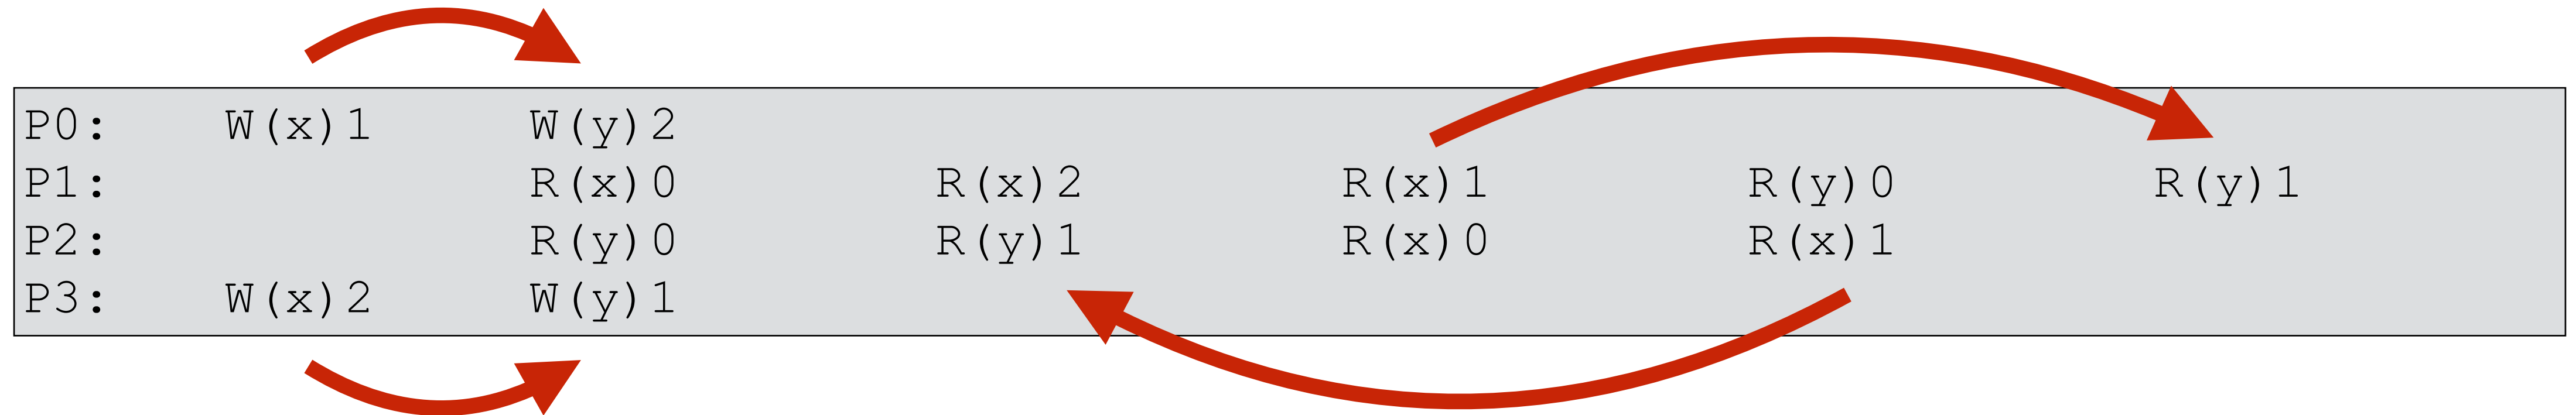
A:	$W_1(y, 2)$	$<$	$R_2(y, 2)$	$<$	$R_2(x, 0)$	$<$	$W_0(x, 1)$	$<$	$R_2(x, 1)$	$<$	$R_1(a, 0)$
B:	$W_1(y, 2)$	$<$	$R_2(y, 2)$	$<$	$R_2(x, 0)$	$<$	$R_1(a, 0)$	$<$	$W_0(x, 1)$	$<$	$R_2(x, 1)$
C:	$W_1(y, 2)$	$<$	$R_2(x, 0)$	$<$	$W_0(x, 1)$	$<$	$R_2(x, 1)$	$<$	$R_1(a, 0)$	$<$	$R_2(y, 2)$
D:	$W_1(y, 2)$	$<$	$W_0(x, 1)$	$<$	$R_2(y, 2)$	$<$	$R_2(x, 0)$	$<$	$R_2(x, 1)$	$<$	$R_1(a, 0)$

Sequentially consistent: A, B

Sequentially inconsistent: C, D

SEQUENTIAL CONSISTENCY: SUFFICIENT CONDITIONS

Sequentially consistent != cache coherent



Sufficient conditions

1. Processors issue memory ops in program order
2. Processors must wait for stores to complete before issuing next memory ops
3. After load, issuing processor waits for load to complete, before issuing next memory operation

Easily implemented with shared (physical) bus

Sufficient, but more than necessary

Example: SGI Origin is SC (MIPS R10000 is dynamically scheduled)

RELAXED CONSISTENCY MODELS

RELAXED CONSISTENCY

Do we really need such a strong model?

Buggy, because of a data race

P0

`x = x + 1`

P1

`x = x + 2`

Unpredictable behavior without appropriate synchronization

So, why to maintain strong consistency anyway?

Compilers neither see multiple processors nor the consistency model

P0

```
A = 1;  
if (B != 0) goto retry;  
// enter critical section
```

P1

```
B = 1;  
if (A != 0) goto retry;  
// enter critical section
```


PROCESSOR CONSISTENCY

Processor consistency (PC): *“writes done by a single processor are received by all other processors in the order in which they were issued, but writes from different processors may be seen in a different order by different processors”*

- Relax write-to-read order (W2R)

- Allow FIFO store buffers, allowing reads to bypass incomplete writes

- W2W, R2R, R2W still ordered

Formal requirement [Goodman 1989]:

- Before load is performed, all previous loads must have been performed

- Before store is performed, all previous memory operations must have been performed

Does not require store atomicity!

Examples: x86, VAX, roughly IBM-370, SUN's Total-Store-Order (includes store atomicity)

See also: PRAM (Pipelined Random Access Memory) consistency

- Also called FIFO consistency

- (Not Parallel Random Access Machine from computability theory)

PROCESSOR CONSISTENCY

Example 1: sequentially inconsistent example

Valid under processor consistency!



Imagine Ps in a chain

Example 2:

Invalid (W2W is preserved)!

Example 3: previous mutual exclusion code?

Unsafe (W2R not preserved)!



```
P0
A = 1;
if (B != 0) goto retry;
// enter critical section
```

```
P1
B = 1;
if (A != 0) goto retry;
// enter critical section
```

RELAXED CONSISTENCY: SAFETY NET


Each relaxed consistency model comes with a safety net

Allowing to revert to sequential consistency

Barriers, syncs, atomics

Puh!

Synchronization accesses vs. ordinary accesses



P0

$x = x + 1$

S

P1

S

$x = x + 2$

Competing: multiple accesses from different processors to the same address, at least one write

PARTIAL STORE ORDER

PC: relax W2R order

SUN's partial store order (PSO): relax W2R and W2W order

Safety net: RMW, memory barriers

Improves store buffer: unordered, coalescing post-retirement store buffer

Much more efficient

WEAK CONSISTENCY

Weak consistency (WC): relax everything (W2W, W2R, R2W, R2R), only separate accesses into sync and non-sync

Properties:

1. Synchronization accesses are sequentially consistent
2. No synchronization access is allowed until all previous writes have completed
3. No access is allowed until all previous synchronization accesses have completed

P0:	W (x) 1	W (x) 2		S	
P1:			R (x) 0	R (x) 2	S R (x) 2
P2:				R (x) 1	S R (x) 2

Key advantage: no need to propagate changes until synchronization occurs

But then, update all memory

Dramatic network traffic reduction

RELEASE CONSISTENCY


Instead of globally updating memory, release consistency (RC) considers locks on areas of memory, and propagates only the locked memory as needed

Definition:

1. Before a non-sync access is performed, all previous acquires by the process must have completed
2. Before a release is performed, all previous reads/writes must have completed
3. Acquire/release is sequentially consistent (RCsc)

Eager: actions for releases

Lazy: for subsequent acquires



RELAXING ALL ORDER - APPROACH 1

Approach 1: using explicit “fence” (memory barrier)

SUN’s Relaxed Memory Order (RMO), Alpha, PowerPC, ARM

x86: “Loads may be reordered with older stores to different locations”

MFENCE, SFENCE, LFENCE

P0:	W (x) 1	W (x) 2			FENCE	
P1:			R (x) 0	R (x) 2	FENCE	R (x) 2
P2:				R (x) 1	FENCE	R (x) 2

RELAXING ALL ORDER - APPROACH 2

Approach 2: annotate loads/stores that do synchronization

Weak consistency (WC): loads/stores can be labeled as “sync”

No reordering allowed across sync operations

Release consistency (RC): specialize loads and stores as acquires and releases

Loads and stores form critical sections, allows reordering into critical sections, but not out

IA64: RCpc (release/acquire obey PC)

Similar to Data-Race-Free-0 (DRF0)

Foundation of C++/Java memory model

P0:	W (x) 1	W (x) 2			FENCE	
P1:			R (x) 0	R (x) 2	FENCE	SR (x) 2
P2:				R (x) 1	FENCE	SR (x) 2

RELAXING ALL ORDER - EXAMPLE

P0

```
A = 1;  
B = 1;  
flag = 1;
```

P1

```
while (!flag);  
print A;  
print B;
```



P0


```
A = 1;  
B = 1;  
SYNCH flag = 1;
```

P1

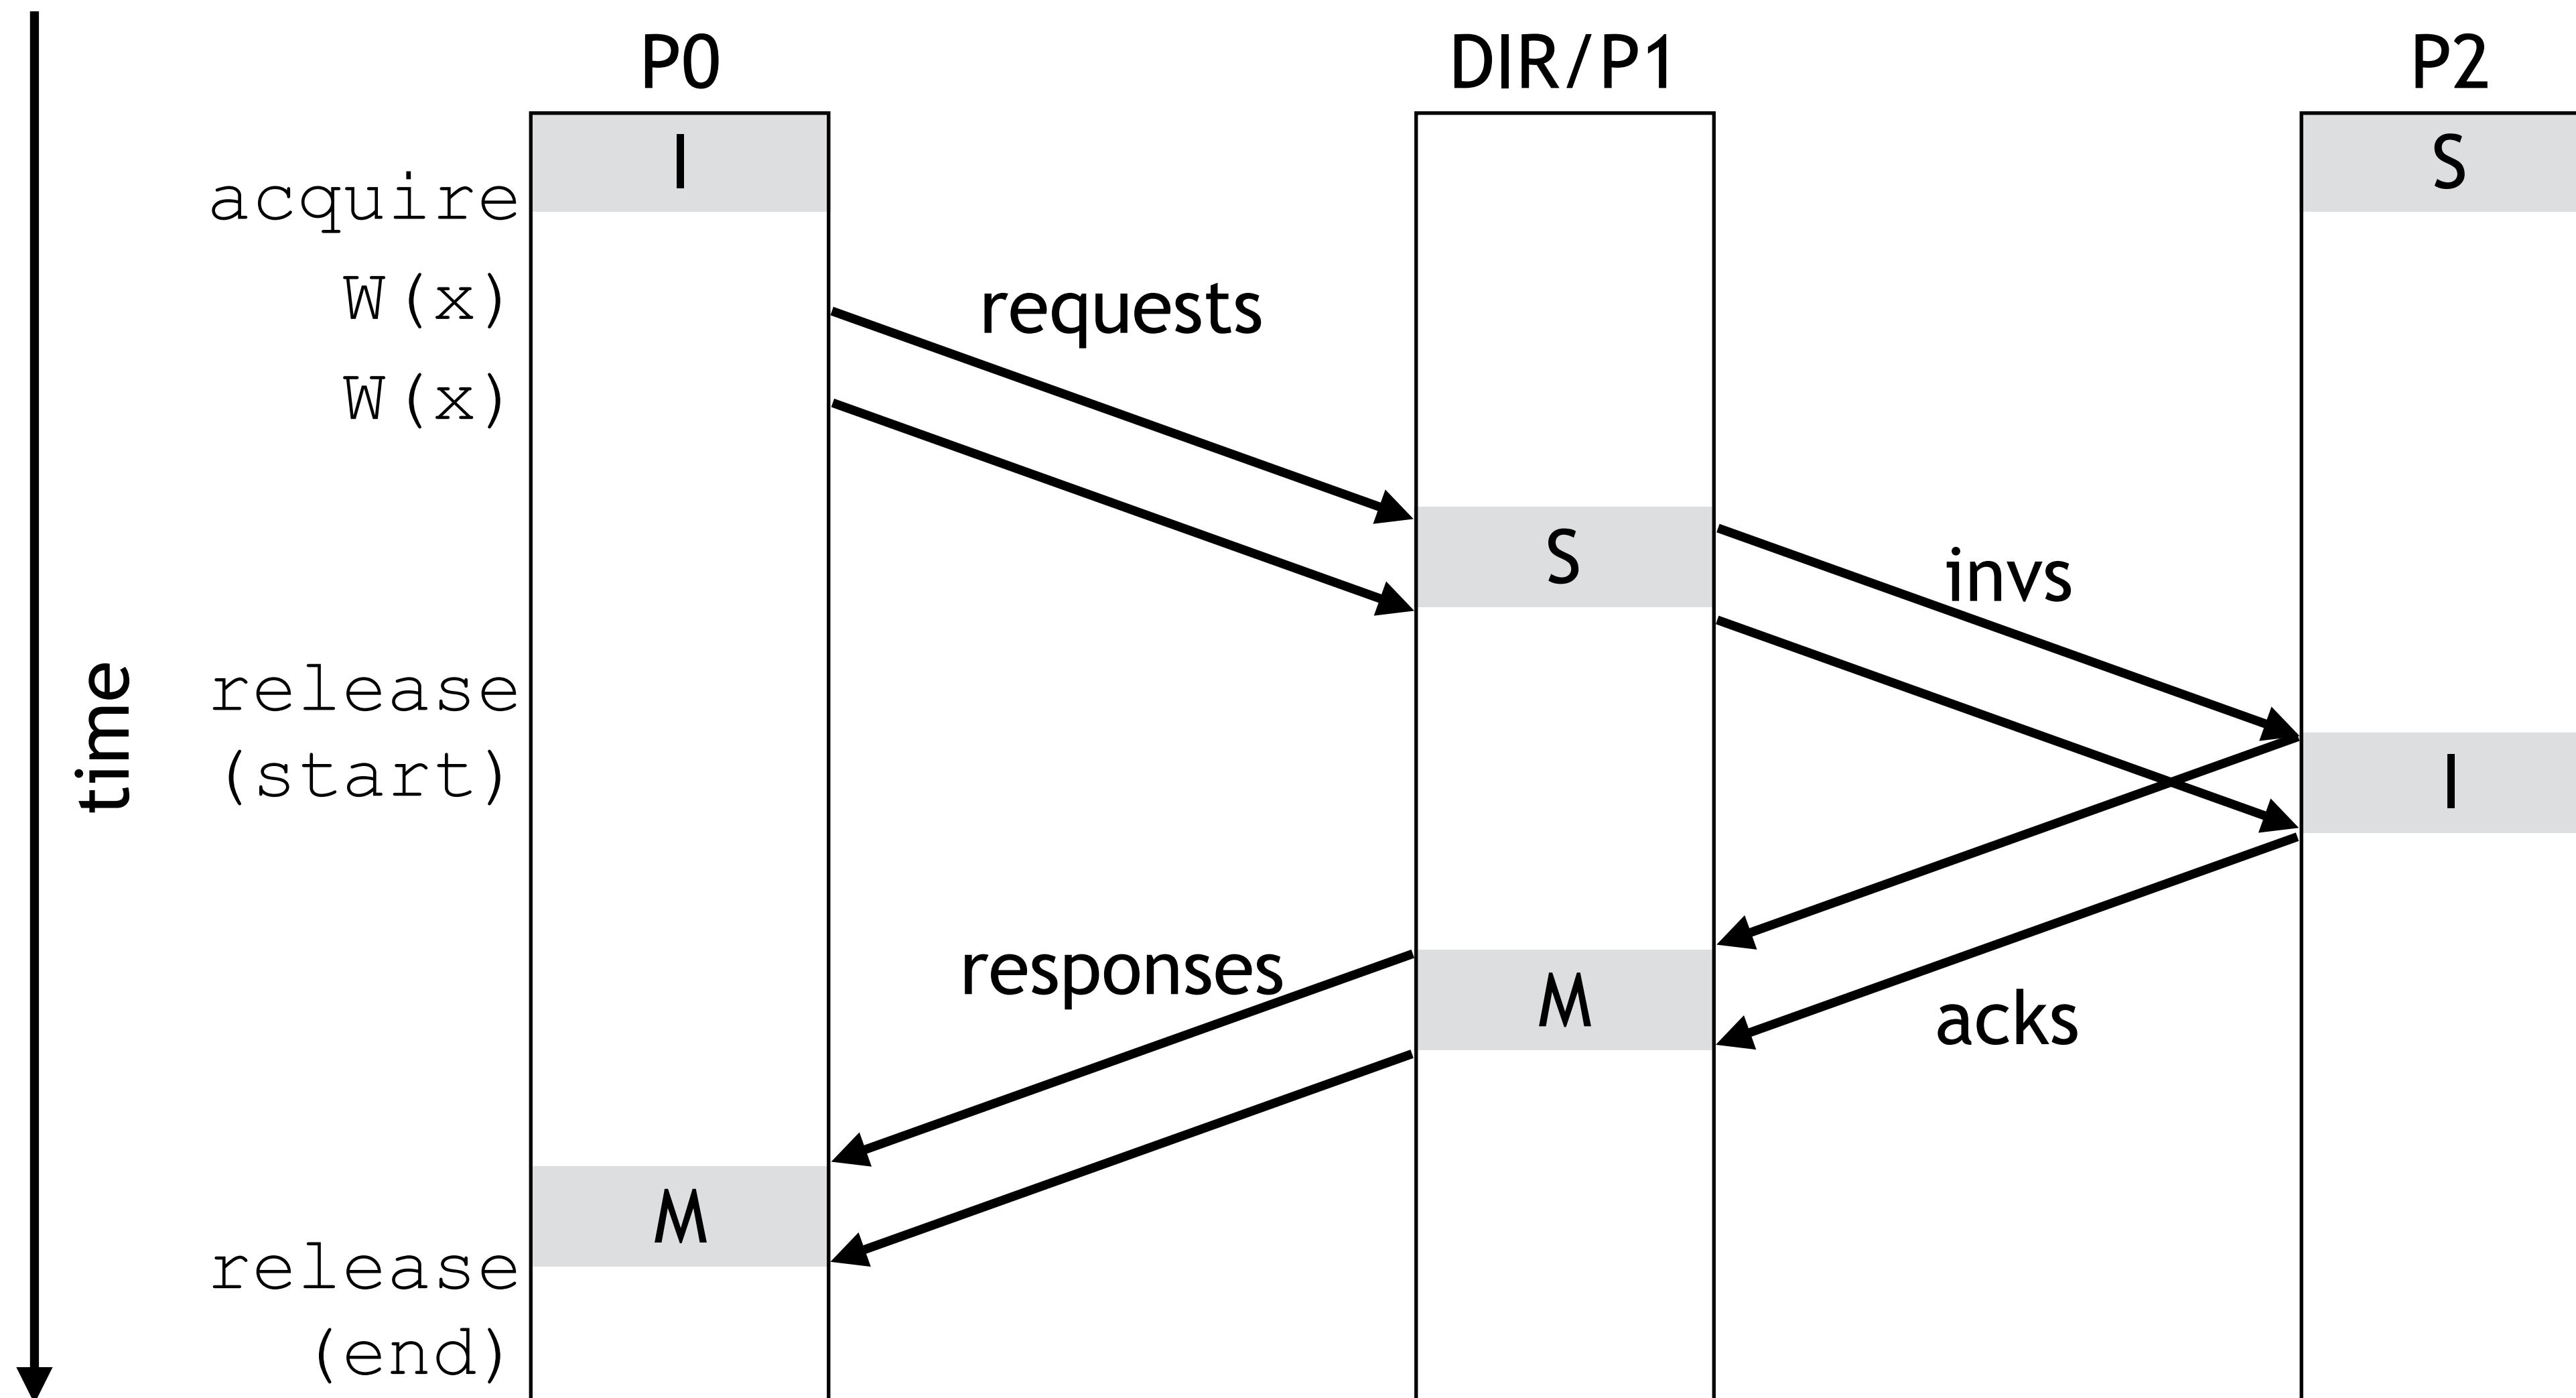
```
while (SYNCH !flag);  
print A;  
print B;
```

COMMENTS

DIRECTORY EXAMPLE WITH SEQUENTIAL CONSISTENCY



DIRECTORY EXAMPLE WITH RELEASE CONSISTENCY



RELAXING ALL ORDER

Compiler still needs to be careful when optimizing:

```
unsigned int i = x;
// x is shared variable

if (i < 2) { // opt: don't copy x, use by ref

    foo: ... // suppose x changes here...

    switch (i) { // opt: implement as jump table
        case 0: ... break;
        case 1: ... break;
        default: ... break;           // opt: range inference
                                      // tells compiler this
                                      // case is impossible,
                                      // drop the check
    }
}
```

SOME COMMENTS

Advanced programming models: Partitioned Global Address Space

UPC: Strict and relaxed operations defined as part of the language

Large-scale DBMS: eventually consistent

Amazon, Google, Facebook, etc

MEMSCALE's On-Demand Consistency Model

Custom, but like weak consistency

Synchronization operations handled in software



	Programmer	Compiler	Hardware	Comment
Strict Consistency	What most/novice programmer expects	Complete disaster!	Global ordering / clock required! No OOO, latency hiding difficult!	Only for uniprocessors
Sequential Consistency	Least astonishing Typically assumed for cache coherence	Disaster! Almost all optimizations are illegal, no reordering!	Disaster! Only one outstanding request! No OOO!	Overkill, most programmers rely on synchronization intrinsics!
Processor Consistency	Sometimes unexpected behavior (membar); however, locks (RMWs) work	May now reorder loads across stores, potential left	Allows for FIFO store buffers & multiple outstanding requests	Typical today x86
Relaxed Consistency (WC,RC,EC)	Very hard (membars where needed)	Sweet, sweet freedom!	Allows for unordered, coalescing SBs & OOO CPUs	

